

Project 2: The NEU 5335 Drone Racing Challenge!



Image credits: [Aerial Sports](#)

"Faster, faster, until the thrill of speed overcomes the fear of death!"

Ladies and gentlemen, boys and girls, children of all ages! Welcome to the 2nd Programming project of CS 5335!! In this project, we'll be flying drones from one place to another and racing through hoops, using path planning with drone dynamics!

Have you ever wondered [how drones fly](#)? How do they hover in the same position in the sky? How do they rotate about their axes and move towards their intended location in the three-dimensional world? We hope you figure all of this as you reach towards the last cells of this notebook!

If you think that Quidditch is the most enticing sport in the world, then you are terribly mistaken. The coolest sport in today's 21st century world is DRONE RACING and if you don't believe this, you should check out this video!



Really cool, right?! But flying with code is even cooler! By the end of the project, you'll also be able to fly just like this!!

Instructions

- Project PA2 will be released on November 11th, Tuesday and will be due on November 24, Monday at 11:59 AM.
- You will not be able to save your work in the release version of the assignment. You **MUST** save a copy to your Drive before you will be able to save any work.
- There are 25 "TODO" items that require you to write code. You will need to complete all of these.
- There are also reflection questions which require a written response. You can find all of these by searching the notebook for "reflection." Your response should go in a text cell directly under each question.
- We may release an extra credit extension to the assignment at a later date

Logistics and Tips

- You need to submit your code and report on gradescope. Please read the submission guidelines at the end of the notebook.
- START EARLY!! This may take some time.
- Press `Shift + Enter` to execute the current code cell, instead of clicking the play button with your mouse.
- Press `b` (b for below) to add a code cell below the current one.
- Press `a` (a for above) to add a code cell above the current one.

Hope you enjoy working on this project as much as we did creating it! Let's [learn to fly!](#) ^_^

▼ Implementation Notes

This assignment uses [GTSAM](#) to represent rotations, translations, etc. GTSAM is primarily a c++ project which offers Python bindings. Most of the available documentation is written for c++, but in general the c++ documentation is applicable to the python bindings.

The drone dynamics used in this assignment are discussed in detail [here](#). This is not required reading, but it may enhance your understanding.

▼ Setup

You'll need to run these cells whenever you fire up collab. But don't worry, it shouldn't take much time!

```
## installation

## imports
from typing import List, Tuple, Dict
import math

# %pip uninstall numpy
%pip install numpy==1.25
import numpy as np
import pandas as pd
import unittest
import plotly.express as px
import plotly.graph_objects as go
%pip install gtsam==4.2
import gtsam

rng = np.random.default_rng(12345)

Requirement already satisfied: numpy==1.25 in /usr/local/lib/python3.11/dist-packages (1.25.0)
Requirement already satisfied: gtsam==4.2 in /usr/local/lib/python3.11/dist-packages (4.2)
Requirement already satisfied: numpy>=1.11.0 in /usr/local/lib/python3.11/dist-packages (from gtsam==4.2) (1.25.0)
Requirement already satisfied: pyparsing>=2.4.2 in /usr/local/lib/python3.11/dist-packages (from gtsam==4.2) (3.2.3)
```

We import `helpers.py` here that'll help us visualize our trees and maps. `gdown` should work fine since it's just one file, but in case you run into any issues, you can also download the `helpers_obstacles.py` file and upload it manually to your notebooks.

```
#downloads the helpers file. This can be commented after running once
!pip install --upgrade --no-cache-dir gdown &> /dev/null
!gdown 1rOEki1n8coTVURLBL7Yh2icbpPc-1fB_

import helpers_obstacles as helpers
# the Drone dynamics model is described in detail in: https://www.roboticsbook.org/S72_drone_actions.html
from helpers_obstacles import Drone, axes, axes_figure

Downloading...
From (original): https://drive.google.com/uc?id=1rOEki1n8coTVURLBL7Yh2icbpPc-1fB
From (redirected): https://drive.google.com/uc?id=1rOEki1n8coTVURLBL7Yh2icbpPc-1fB&confirm=t&uuid=f2b6493e-d3fe-4f8d-bfc7-a66c1
To: /content/helpers_obstacles.py
100% 61.3k/61.3k [00:00<00:00, 4.80MB/s]
```

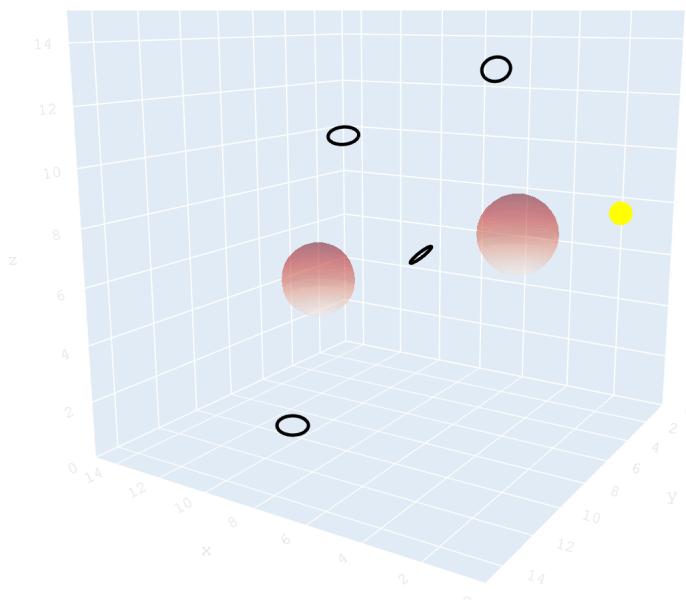
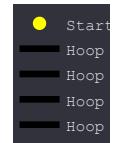
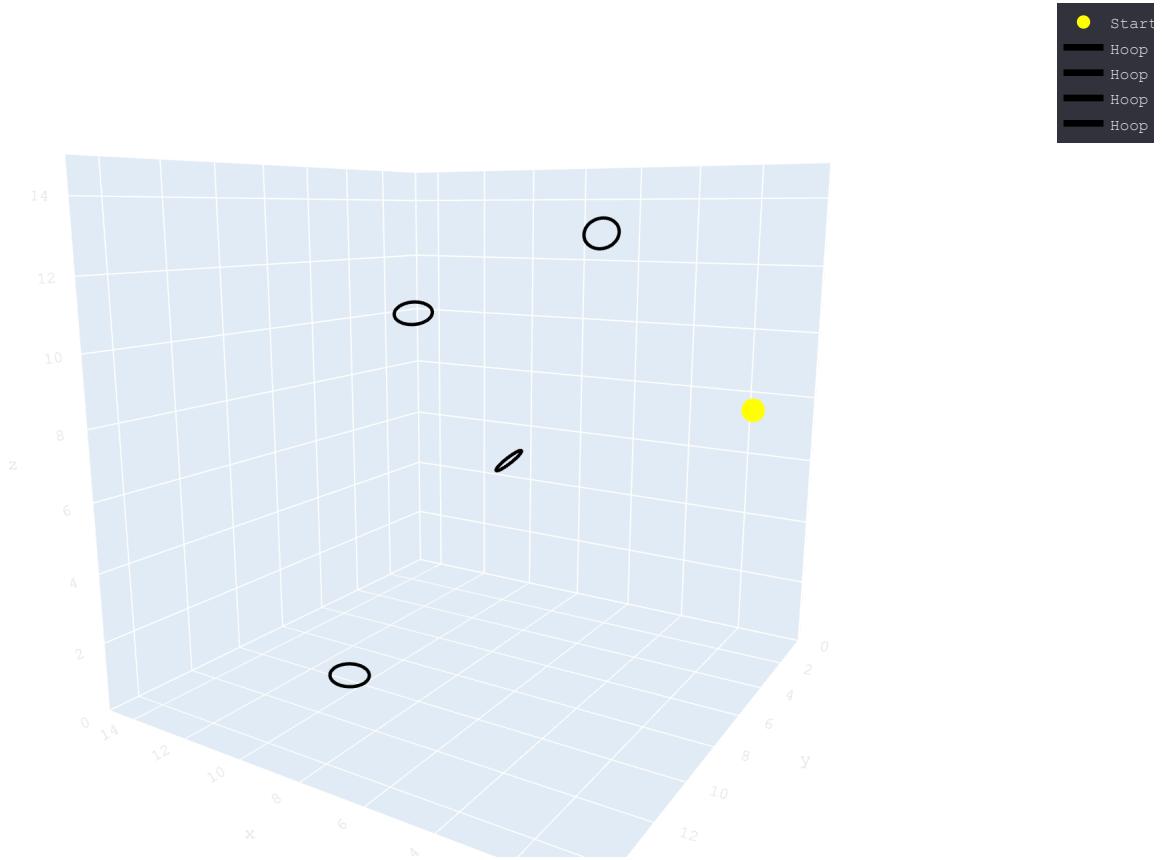
▼ Part 1: Rapidly-exploring Random Trees (RRT)! In 3 Dimensions!

- RRT again? Yes you read that right! But this time, we're taking it up a notch! 😊
- Given the coordinates of our starting location and the destination location, how does a drone fly from the start to the goal? One easy yet effective way is to execute the RRT algorithm for the drone!
- In this section, we shall implement the RRT algorithm in three dimensions (3D), because well, we live in a 3 dimensional world.
- In the previous assignment, each node in our RRT tree was a (x, y) tuple. Now that we're in three dimensions, we require three coordinates (x, y, z) to represent the position of the drone. We shall use `gtsam.Point3` to represent this.

The NEU Drone Racing Track!

Let's first quickly load up the racing track. Take a look at those suspended HOOPS!!! In mid air!! We're gonna do RRT and fly through 'em!

```
start_race = gtsam.Pose3(r=gtsam.Rot3.Yaw(math.radians(45)), t=gtsam.Point3(1, 3, 8))
helpers.drone_racing_path(helpers.get_hoops(), start_race, [])
helpers.drone_racing_path_with_obstacles(helpers.get_hoops(), start_race, [], obstacles=helpers.get_obstacles_easy())
```



- We now code up the helper functions that'll help us execute the RRT algorithm. Please read the docstrings carefully for each of them, the arguments to the function, and the return types.
- You can refer to [Section 5.5.6](#) in the textbook for some helpful insights, especially for TODO 3.

```
# TODO 1
def generate_random_point(target: gtsam.Point3) -> gtsam.Point3:
    """
    This function generates a random node in the 3 dimensional configuration space of (10x10x10) and returns it.
    You must ensure that there is atleast a 20% chance with which the target node itself is returned.

    Hints:
    - Use rng for random number generation

    Arguments:
    - target: gtsam.Point3 (the goal itself!)

    Returns:
    - node: gtsam.Point3 - a random configuration of (x, y, z) which should be within the (10x10x10) space
    """

    node = None

    ##### Student code here #####
    if rng.random() < 0.2: # 20% chance
        return target

    # Otherwise, picking a random point in the 10x10x10 space
    x = rng.uniform(0, 10)
    y = rng.uniform(0, 10)
    z = rng.uniform(0, 10)

    node = gtsam.Point3(x, y, z)

    # raise NotImplementedError("generate_random_point is not implemented") # Commenting out since you're implementing

    ##### End student code #####
    return node
```

```
# TODO 2
def distance_euclidean(point1: gtsam.Point3, point2: gtsam.Point3) -> float:
    """
    This function computes the euclidean distance between two 3-D points.

    Hints:
    - Use np.linalg.norm to compute the norm

    Arguments:
    - point1: gtsam.Point3
    - point2: gtsam.Point3

    Returns:
    - distance_euclidean: float
    """

    distance_euclidean = None

    ##### Student code here #####
    # Just the straight-line distance formula in 3D
    # Like measuring with a ruler between two points

    # Converting to numpy arrays for easier math
    p1 = np.array([point1[0], point1[1], point1[2]])
    p2 = np.array([point2[0], point2[1], point2[2]])

    # Calculating the difference and find its length
    distance_euclidean = np.linalg.norm(p1 - p2)
    # raise NotImplementedError("distance_euclidean is not implemented") # Commenting out as I am providing implementation

    ##### End student code #####
    return distance_euclidean
```

- Once we've sampled a random node, we need to find a potential parent for that node. We can do this by finding that node in the tree which is closest to the newly sampled node.
- A naive way of doing this is to loop through all the nodes in the tree, compute the distance and pick the node with the minimum distance.
- However, for this function, you must use numpy vectorization and parallelize your implementation. A good resource to learn this: [Link](#)

```
# TODO 3
def find_nearest_node(rrt: List[gtsam.Point3], node: gtsam.Point3):
    """
    Given the current RRT tree and the newly sampled node, this function returns the node in the tree which is CLOSEST
    to the newly sampled node, as well as the index of that node.

    This can be done naively by looping through each node and computing the distance, but you need to parallelize it!

    Hints:
    - Refer to the textbook on how to vectorize it

    Arguments:
    - rrt: List[gtsam.Point3] (the current tree)
    - node: gtsam.Point3 (the newly sampled point)

    Returns:
    - nearest_node: gtsam.Point3 (the nearest node!)
    - index: int (the index of the nearest node)
    """

    nearest_node = None
    index = None

    ##### Student code here #####
    # Stacking all the tree nodes into a matrix (each row is a node)
    tree_array = np.array([[p[0], p[1], p[2]] for p in rrt])

    # target node as an array
    node_array = np.array([node[0], node[1], node[2]])

    # Calculating distances to all nodes at once (broadcasting magic!)
    # subtracting the node from every tree point, then finding the length
    distances = np.linalg.norm(tree_array - node_array, axis=1)

    # Finding the index with minimum distance
    index = np.argmin(distances)
    nearest_node = rrt[index]

    #raise NotImplementedError("find_nearest_node is not implemented")

    ##### End student code #####
    return nearest_node, index
```

- Once we've sampled a node and found the node closest to it in the RRT tree, we need to find a node which would take us **in the direction** of the sampled node.
- In this 3-dimensional scenario, we shall call this the `steer_node` as we wish to steer our drone in the direction of the newly sampled node.
- We adopt a simple and naive strategy: we turn towards the target, and drive some **fraction of the distance**. (Hint: refer to the textbook!)

```
# TODO 4
def steer_naive(parent: gtsam.Point3, target: gtsam.Point3, fraction = 0.2):
    """
    This function steers the drone towards the target point, going a fraction of the displacement
    It returns the new 'steer_node' which takes us closer to the destination.

    Arguments:
    - parent: gtsam.Point3
    - target: gtsam.Point3

    Returns:
    - steer_node: gtsam.Point3
    """

    # Implementation goes here
```

```

steer_node = None

##### Student code here #####
# Converting to numpy for easier vector math
parent_array = np.array([parent[0], parent[1], parent[2]])
target_array = np.array([target[0], target[1], target[2]])

# Moving from parent toward target by 'fraction' of the distance
steer_array = parent_array + fraction * (target_array - parent_array)

# Converting back to Point3
steer_node = gtsam.Point3(steer_array[0], steer_array[1], steer_array[2])

# raise NotImplementedError("steer_naive is not implemented")

##### End student code #####
return steer_node

```

```

class TestRRT(unittest.TestCase):
    def test_generate_random_point(self):
        for _ in range(5):
            node = generate_random_point(gtsam.Point3(4,5,6))
            assert 0 <= node[0] <= 10
            assert 0 <= node[1] <= 10
            assert 0 <= node[2] <= 10

    def test_distance_euclidean(self):
        pt1 = gtsam.Point3(2.70109492, 4.55796488, 2.93292049)
        pt2 = gtsam.Point3(4, 7, 2)
        self.assertAlmostEqual(distance_euclidean(pt1, pt2), 2.9190804346571446, 2)

    def test_find_nearest_node(self):
        pt1 = gtsam.Point3(1,2,3)
        pt2 = gtsam.Point3(0.90320894, 3.55218386, 3.71979848)
        pt3 = gtsam.Point3(1.52256715, 4.24174709, 3.37583879)
        pt4 = gtsam.Point3(1.56803165, 4.10257537, 2.795647)
        pt5 = gtsam.Point3(2.68087164, 3.63713802, 4.25464017)
        new_point = gtsam.Point3(3.74935314, 3.2575652 , 5.20840562)
        rrt = [pt1, pt2, pt3, pt4, pt5]
        answer, index = find_nearest_node(rrt, new_point)
        assert (answer==pt5).all()

    def test_steer_naive(self):
        pt1 = gtsam.Point3(3.80319106, 2.49123788, 2.60348781)
        pt2 = gtsam.Point3(3.81712339, 0.33173367, 0.51835128)
        answer = gtsam.Point3(3.80597753, 2.05933704, 2.1864605)
        steer_node = steer_naive(pt1, pt2)
        assert(np.allclose(answer, steer_node, atol=1e-2))

    suite = unittest.TestSuite()
    suite.addTest(TestRRT('test_generate_random_point'))
    suite.addTest(TestRRT('test_distance_euclidean'))
    suite.addTest(TestRRT('test_find_nearest_node'))
    suite.addTest(TestRRT('test_steer_naive'))

    unittest.TextTestRunner().run(suite)

...
-----
Ran 4 tests in 0.016s

OK
<unittest.runner.TextTestResult run=4 errors=0 failures=0>

```

Putting it all together!

Let's now use all the functions coded up above and write the RRT loop. Here's the outline of the algorithm for your reference:

1. Start with the RRT tree containing the start node. Our aim is to grow this tree with every iteration of the loop.
2. Sample out a random node in the configuration space. Make sure you return the target node with a 20% probability.
3. Find the node nearest to the newly sampled node in the current tree, and make sure you keep track of the parent node. For every node that we add to the tree, you also need to store the index of its parent.

4. Find the "steer node" - a node in the direction of the sampled node, and add it to the RRT tree. No need to check for obstacles, we don't have any! 😊
5. Repeat steps 2 and 3 until the distance of the latest node in the tree and the target node is less than the threshold. As soon as this terminating condition is reached, you can return the tree and the parents list.

This function is the generalized RRT function which we shall be utilizing multiple times in this project. **It uses other helper functions as arguments!! So make sure you use the exact names as given in the function header.** This is **not** limited to the functions we've coded above.

```
# TODO 5
def run_rrt(start, target, generate_random_node, steer, distance, find_nearest_node, threshold):
    ...
    This function is the main RRT loop and executes the entire RRT algorithm.
    Follow the steps outlined above. You should keep sampling nodes until the terminating condition is met.

    Please use the same function names as given in the function definition.

    Arguments:
        - start: the start node, it could be gtsam.Point3 or gtsam.Pose3.
        - target: the destination node, it could be gtsam.Point3 or gtsam.Pose3.
        - generate_random_node: this function helps us randomly sample a node
        - steer: this function finds the steer node, which takes us closer to our destination
        - distance: this function computes the distance between the two nodes in the tree
        - find_nearest_node: this function finds the nearest node to the randomly sampled node in the tree
        - threshold: float, this is used for the terminating the algorithm

    Returns:
        - rrt: List[gtsam.Point3] or List[gtsam.Pose3], contains the entire tree
        - parents: List[int], contains the index of the parent for each node in the tree
        ...
        ...

    rrt = []
    parents = []
    max_iterations = 2000
    rrt.append(start)
    parents.append(-1)

    for i in range(max_iterations):
        ##### Student code here #####
        # Step 1: Sampling a random node (with goal bias)
        random_node = generate_random_node(target)

        # Step 2: Finding the closest node in our tree
        nearest_node, nearest_idx = find_nearest_node(rrt, random_node)

        # Step 3: Steering toward the random node
        new_node = steer(nearest_node, random_node)

        # Step 4: Adding this new node to our tree
        rrt.append(new_node)
        parents.append(nearest_idx)

        # Step 5: Checking if we're close enough to the goal
        dist_to_goal = distance(new_node, target)
        if dist_to_goal < threshold:
            # Success! reached the goal
            return rrt, parents

    ##### End student code #####
    # If max_iterations reached without finding a path within the threshold
    return rrt, parents
```

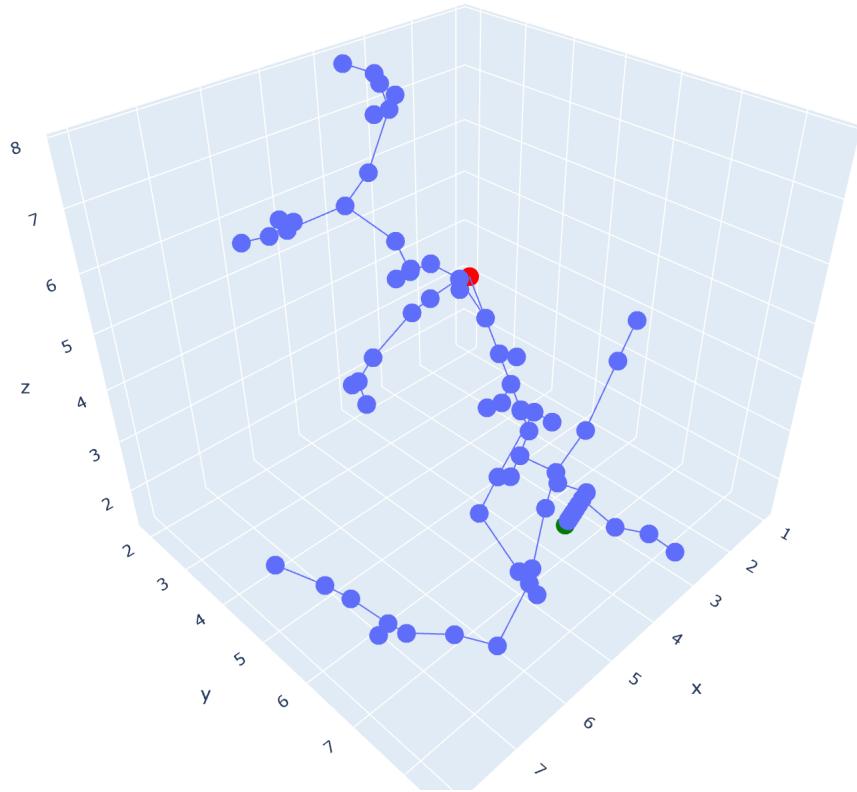
▼ Visualizing the RRT Tree and Path!

Alright, with this, we've completed an implementation of the RRT algorithm in 3 dimensions! Let's see what our tree and path look like!

```
start_rrt_3d = gtsam.Point3(1,2,3)
target_rrt_3d = gtsam.Point3(4,7,2)
rrt_3d, parents_rrt_3d = run_rrt(start_rrt_3d, target_rrt_3d, generate_random_point, steer_naive, distance_euclidean, find_nearest_node)
print("Nodes in RRT: ", len(rrt_3d))
```

Nodes in RRT: 73

```
helpers.visualize_tree(rrt_3d, parents_rrt_3d, start_rrt_3d, target_rrt_3d)
```



- Remember we kept track of the parents of every node added to the RRT tree? We shall use this information to print out the path from the start node to the target node.
- You need to backtrack and store the node's parent in a `path` list. We've already implemented this function for you 😊

```
def get_rrt_path(rrt: List[gtsam.Pose3], parents: List[int]) -> List[gtsam.Pose3]:
    path = []
    i = len(rrt) - 1
    path.append(rrt[i])

    while(parents[i] != -1):
        next = rrt[parents[i]]
        path.append(next)
        i = parents[i]

    path.reverse()
    return path
```

Reflection Questions:

- If you look closely at the tree and the path, you'll notice that as we reach close to the goal (the green node), the nodes are extremely close to one another! We begin to take really small steps and we converge very slowly towards it. Can you explain why this is happening?

▼ Response:

It's the fraction parameter in `steer_naive`, namely, at each step we only move 20% of the distance. This works fine when we are far away but close to the goal everything becomes tiny, such that if the nearest node is 0.5m from goal and we sample nearby, 20% of that is only 0.1m, and hence we end up taking baby steps near the goal. That part of the tree gets super crowded because we just continually add nodes very near each other. Doesn't help that we have 20% goal bias either, so it constantly samples near target. Pretty much creates this convergence zone where progress is a crawl, would probably work way better if it was fixed step size instead of fraction-based.

▼ Part 2: Drone Dynamics

Now that we have a taste for path finding in 3D, let's see how a drone actually flies before we start steering it.

The kinematics of quadrotors and multicopters are those of simple rigid 3D bodies. The kinematics equations are most useful for navigation and control when expressed in the navigation frame \mathcal{N} , which for MAV applications is almost universally assumed to be non-rotating and aligned with gravity. In this project, we use the **ENU (East - North - Up)** navigation frame. This means that the x-axis points towards the east, the y-axis points towards the north and the z-axis points upwards.

We also define a body frame \mathcal{B} as having its origin at the center of mass of the vehicle. Following convention in aerospace applications, we fix the the x-axis as pointing to the front of the vehicle (not always the direction of travel), the y-axis as pointing to the left, and the z-axis pointing up, the so-called **FLU (Forward - Left - Up)** convention.

We then define, respectively,

- the vehicle's position $r^n \doteq [x, y, z]^T$,
- its linear velocity $v^n = \dot{r}^n \doteq [u, v, w]^T$,
- the attitude $R_b^n \doteq [i^b, j^b, k^b] \in SO(3)$, a 3×3 rotation matrix from \mathcal{B} to \mathcal{N}

Above the superscript n and b denote quantities expressed in the *navigation* and *body* frame, respectively.

As we're now using the dynamics of the drone, we also need to represent the orientation of the drone as a rotation matrix. As you may have already guessed, we have an elegant structure that meets our needs - `gtsam.Pose3`.

The attitude and position can be represented by `gtsam.Rot3` and `gtsam.Point3` objects respectively. And they can be combined together to represent the *pose* of the drone using `gtsam.Pose3` objects.

For example, if the current position of the drone is at the origin of \mathcal{N} and \mathcal{B} aligns with \mathcal{N} , the attitude is an identity matrix and the position is $[0, 0, 0]$. Let's see how we can represent this in code.

```
position = gtsam.Point3(0, 0, 0)
attitude = gtsam.Rot3()
pose = gtsam.Pose3(r = attitude, t = position)

print(f"Position: {position}")
print(f"Attitude: {attitude}")
print(f"Pose: {pose}")

Position: [0. 0. 0.]
Attitude: R: [
    1, 0, 0;
    0, 1, 0;
    0, 0, 1
]
Pose: R: [
    1, 0, 0;
    0, 1, 0;
    0, 0, 1
]
t: 0 0 0
```

We're going to use the following test suite to check your implementation of the different methods in this part. Feel free to add more test cases in the suite for your own testing.

```
class TestDroneDynamics(unittest.TestCase):
    def test_compute_attitude_from_ypqr(self):
        yaw = math.radians(45)
        pitch = math.radians(30)
        roll = math.radians(60)
```

```

expected_attitude = gtsam.Rot3(
    [0.612372, 0.612372, -0.5],
    [-0.0473672, 0.65974, 0.75],
    [0.789149, -0.435596, 0.433013]
)
actual_attitude = compute_attitude_from_ypr(yaw, pitch, roll)

assert(actual_attitude.equals(expected_attitude, tol=1e-2))

def test_compute_force(self):
    attitude = gtsam.Rot3(
        [0.612372, 0.612372, -0.5],
        [-0.0473672, 0.65974, 0.75],
        [0.789149, -0.435596, 0.433013]
    )
    thrust = 20.0

    expected_force = gtsam.Point3(15.78, -8.71, 8.66)
    actual_force = compute_force(attitude, thrust)

    assert(np.allclose(actual_force, expected_force, atol=1e-2))

def test_compute_terminal_velocity(self):
    force = gtsam.Point3(15.78, -8.71, 8.66)

    expected_terminal_velocity = gtsam.Point3(19.27, -14.32, 14.27)
    actual_terminal_velocity = compute_terminal_velocity(force)

    assert(np.allclose(actual_terminal_velocity, expected_terminal_velocity, atol=1e-2))

```

The attitude of the drone can also be defined using the yaw, pitch and roll angles of the drone with respect to \mathcal{N} . Our first task now will be to compute the attitude of the drone, R_b^n , given the yaw, pitch and roll angles of the drone.

```

# TODO 6
def compute_attitude_from_ypr(yaw: float, pitch: float, roll: float) -> gtsam.Rot3:
    ...
    Uses yaw, pitch and roll angles to compute the attitude of the drone

    Arguments:
        - yaw: float (in radians)
        - pitch: float (in radians)
        - roll: float (in radians)

    Hint: Use help(gtsam.Rot3) to see different constructors

    Returns:
        - attitude: gtsam.Rot3
        ...
        attitude = None

        ##### Student code here #####
        attitude = gtsam.Rot3.Ypr(yaw, pitch, roll)

    # raise NotImplementedError("compute_attitude_from_ypr is not implemented")

    ##### End student code #####
    return attitude

```

```

suite = unittest.TestSuite()
suite.addTest(TestDroneDynamics('test_compute_attitude_from_ypr'))

unittest.TextTestRunner().run(suite)

.
-----
Ran 1 test in 0.009s

OK
<unittest.runner.TextTestResult run=1 errors=0 failures=0>

```

Let us take the [Skydio 2+](#) as a prototypical and relevant example. Its mass, with battery, is $800g$, but let's just say $1kg$ to make the calculations easier. Its dimensions are $229mm \times 274mm$, but let's assume the motors are at $(\pm 0.10m, \pm 0.10m)$, which is not too far off.

✓ Flying the drone

The four rotors on the drone provide a thrust upwards (in the body frame) to fly the drone. In general, the force F_z^b aligned with the body z-axis will be:

$$F_z^b = \sum_{i=1}^4 f_i$$

where f_i is the force applied by each rotor.

Of course, when we *tilt* the quadrotor forwards, we will direct some of that thrust towards generating horizontal acceleration. To get a handle on this, we need to calculate the thrust in the *navigation* frame, but this is just a matter of multiplying with the attitude R_b^n :

$$F^n = R_b^n \begin{bmatrix} 0 \\ 0 \\ F_z^b \end{bmatrix} = \hat{z}_b^n F_z^b$$

Our next task is to compute the force vector F^n , given attitude R_b^n and upwards thrust in the body frame, F_z^b .

```
# TODO 7
def compute_force(attitude: gtsam.Rot3, thrust: float) -> gtsam.Point3:
    """
    Computes the force vector given attitude and thrust in the body frame

    Arguments:
        - attitude: gtsam.Rot3, nRb for the drone
        - thrust: float, the upwards thrust produced by the 4 rotors

    Returns:
        - force: gtsam.Point3, the resultant force vector
    """

    force = None

    ##### Student code here #####
    # Thrust vector in body frame (pointing up in FLU convention)
    thrust_body = np.array([0, 0, thrust])

    # Rotating to navigation frame using the attitude matrix
    force_nav = attitude.matrix() @ thrust_body

    force = gtsam.Point3(force_nav[0], force_nav[1], force_nav[2])

    #raise NotImplementedError("compute_force is not implemented")

    ##### End student code #####
    return force
```

```
suite = unittest.TestSuite()
suite.addTest(TestDroneDynamics('test_compute_force'))

unittest.TextTestRunner().run(suite)

.
-----
Ran 1 test in 0.008s

OK
<unittest.runner.TextTestResult run=1 errors=0 failures=0>
```

✓ Drag and Terminal Velocity

Constant forward thrust, as calculated above, does *not* mean that the drone will continue accelerating, because of *drag*. In fact, from the spec-sheet of the Skydio-2 we see that the top speed (in autonomous mode) is 36 Mph , which is about 16 m/s , and the theoretical top

speed is probably more like 20 m/s . Drag force increases *quadratically* with velocity:

$$F_{\text{drag}} \propto v^2$$

$$F_{\text{drag}} = k_d v^2$$

where k_d is the drag coefficient. Using the specs from Skydio-2, we compute this $k_d = 0.0425$.

This drag force has to exactly balance the maximum forward thrust at terminal velocity:

$$F^n = k_d v_{\text{terminal}}^2$$

$$v_{\text{terminal}}^2 = \frac{F^n}{k_d} \implies v_{\text{terminal}} = \sqrt{\frac{F^n}{k_d}}$$

Our task now is to calculate the terminal velocity vector, v_{terminal} , given the force vector, F^n and the drag coefficient, k_d .

```
# TODO 8
def compute_terminal_velocity(force: gtsam.Point3, kd: float = 0.0425) -> gtsam.Point3:
    ...
    Uses the force vector and drag coefficient to compute the terminal velocity of the drone

    Arguments:
        - force: gtsam.Point3, the force vector in the navigation frame
        - kd: float, drag coefficient

    Returns:
        - terminal_velocity: gtsam.Point3, the maximum velocity vector
        ...

    terminal_velocity = None

    # Adding small epsilon in computating force for numerical stability (avoid sqrt(0))
    eps = 1e-6

    ##### Student code here #####
    # Calculating terminal velocity for each component
    # adding eps to avoid sqrt(0) which can cause numerical issues
    vx = np.sign(force[0]) * np.sqrt(abs(force[0]) / (kd + eps))
    vy = np.sign(force[1]) * np.sqrt(abs(force[1]) / (kd + eps))
    vz = np.sign(force[2]) * np.sqrt(abs(force[2]) / (kd + eps))

    terminal_velocity = gtsam.Point3(vx, vy, vz)

    #raise NotImplementedError("compute_terminal_velocity is not implemented")

    ##### End student code #####
    return terminal_velocity
```

```
suite = unittest.TestSuite()
suite.addTest(TestDroneDynamics('test_compute_terminal_velocity'))

unittest.TextTestRunner().run(suite)

.
-----
Ran 1 test in 0.004s

OK
<unittest.runner.TextTestResult run=1 errors=0 failures=0>
```

Part 3: Steering with Terminal Velocity

Now that we know how the thrust applied by the rotors flies the drone with a certain velocity, we can use that to explore the RRT tree.

However, since we learned that the drones now have *poses* instead of just positions, we will need to update our random node generating function to return random poses. Our first task, in this section, is to return a random pose node with the same random position component as before, but a new random attitude component using a random yaw, pitch and roll $\in [-60^\circ, 60^\circ]$.

Again, we're going to use the following test suite to check your implementation of the different methods in this part. Feel free to add more test cases in the suite for your own testing.

```

class TestSteeringWithTerminalVelocity(unittest.TestCase):
    def test_generate_random_pose(self):
        target_node = gtsam.Pose3(r = gtsam.Rot3.Yaw(math.radians(45)), t = gtsam.Point3(8, 5, 6))
        for _ in range(5):
            random_node = generate_random_pose(target_node)
            assert(np.all(np.greater_equal(random_node.translation(), gtsam.Point3(0, 0, 0))))
            assert(np.all(np.less_equal(random_node.translation(), gtsam.Point3(10, 10, 10))))
            assert(np.all(np.greater_equal(random_node.rotation().ypr(), gtsam.Point3(math.radians(-60), math.radians(-60), math.radians(-60)))))
            assert(np.all(np.less_equal(random_node.rotation().ypr(), gtsam.Point3(math.radians(60), math.radians(60), math.radians(60)))))

    def test_find_nearest_pose(self):
        rrt_tree = [gtsam.Pose3(
            r=gtsam.Rot3([1, 0, 0],
                         [0, 1, 0],
                         [0, 0, 1]),
            t=gtsam.Point3(1, 2, 3)),
        gtsam.Pose3(
            r=gtsam.Rot3([0.771517, -0.617213, 0],
                         [0.0952381, 0.119048, -0.97619],
                         [0.617213, 0.771517, 0.154303]),
            t=gtsam.Point3(2.70427, 3.90543, 3.85213)),
        gtsam.Pose3(
            r=gtsam.Rot3([0.601649, -0.541882, 0.302815],
                         [-0.301782, -0.62385, -0.516772],
                         [0.627501, 0.29376, -0.721074]),
            t=gtsam.Point3(4.42268, 5.08119, 2.01005)),
        gtsam.Pose3(
            r=gtsam.Rot3([-0.696943, 0.589581, -0.36631],
                         [-0.664345, -0.416218, 0.594076],
                         [0.204431, 0.679463, 0.704654]),
            t=gtsam.Point3(5.40351, 6.86933, 3.83104)),
        gtsam.Pose3(
            r=gtsam.Rot3([-0.0686996, 0.218721, -0.818805],
                         [-0.796488, -0.297401, -0.0126152],
                         [-0.340626, 0.900832, 0.269211]),
            t=gtsam.Point3(1.43819, 5.96437, 4.97769))]
        new_node = gtsam.Pose3(
            r=gtsam.Rot3([0.682707, 0.661423, 0.310534],
                         [-0.626039, 0.748636, -0.218217],
                         [-0.376811, -0.0454286, 0.925176]),
            t=gtsam.Point3(5.65333, 5.65964, 1.60624))
        expected_nearest_node = rrt_tree[2]
        expected_index = 2
        actual_nearest_node, actual_index = find_nearest_pose(rrt_tree, new_node)
        assert(actual_nearest_node.equals(expected_nearest_node, tol=1e-1))
        assert(actual_index == expected_index)

    def test_steer_with_terminal_velocity(self):
        current_node = gtsam.Pose3(gtsam.Rot3.Yaw(math.radians(90)), gtsam.Point3(1, 2, 3))
        new_node = gtsam.Pose3(gtsam.Rot3.Pitch(math.radians(45)), gtsam.Point3(8, 5, 6))

        expected_steer_node = gtsam.Pose3(gtsam.Rot3(
            [0.37, -0.86, 0],
            [0.31, 0.13, -0.87],
            [0.86, 0.37, 0.37])
        ), gtsam.Point3(3.00, 3.31, 4.31))
        actual_steer_node = steer_with_terminal_velocity(current_node, new_node)

        assert(actual_steer_node.equals(expected_steer_node, tol=1e-1))

```

TODO 9

def generate_random_pose(target: gtsam.Pose3) -> gtsam.Pose3:

'''

This function generates a random node in the pose configuration space (10x10x10) and returns it.
 You must ensure that there is atleast a 20% chance with which the target node itself is returned.
 The attitude can be randomly sampled via yaw, pitch and roll angles between -60 to 60 degrees.

Hints:

- Use rng for random number generation
- Use compute_attitude_from_ypr function

Arguments:

- target: gtsam.Pose3, the target pose for the RRT

Returns:

```

- node: gtsam.Pose3, random pose or target pose
...

node = None

##### Student code here #####
if rng.random() < 0.2: # 20% chance return target
    return target

# Random position
x = rng.uniform(0, 10)
y = rng.uniform(0, 10)
z = rng.uniform(0, 10)

# Random orientation (yaw, pitch, roll) between -60 and 60 degrees
yaw = rng.uniform(-60, 60) * np.pi / 180
pitch = rng.uniform(-60, 60) * np.pi / 180
roll = rng.uniform(-60, 60) * np.pi / 180

# Creating the pose with random position and orientation
attitude = compute_attitude_from_ypr(yaw, pitch, roll)
position = gtsam.Point3(x, y, z)
node = gtsam.Pose3(attitude, position)

# raise NotImplementedError("generate_random_pose is not implemented")

##### End student code #####
return node

```

```

suite = unittest.TestSuite()
suite.addTest(TestSteeringWithTerminalVelocity('test_generate_random_pose'))

unittest.TextTestRunner().run(suite)

```

```

.
-----
Ran 1 test in 0.007s

OK
<unittest.runner.TextTestResult run=1 errors=0 failures=0>

```

Since, we are using poses now, we will also need to update the way we have been calculating the distances between two nodes. You do not need to implement this, you can use the distance function we have provided in `helpers` as `distance_between_poses()`. Run `help(helpers.distance_between_poses)` to know how to use this function.

```

help(helpers.distance_between_poses)

Help on function distance_between_poses in module helpers_obstacles:

distance_between_poses(pose1: gtsam.gtsam.Pose3, pose2: gtsam.gtsam.Pose3) -> float

```

We will now implement the function to find the nearest node in the RRT tree again, using the new distance function. This one is harder to vectorize, so we can implement it in the naive way using a `for` loop.

```

# TODO 10
def find_nearest_pose(rrt: List[gtsam.Pose3], node: gtsam.Pose3):
    ...
    This function finds the nearest node in the current RRT tree to the newly sampled node.

    Arguments:
    - rrt: List[gtsam.Pose3] (a list of nodes currently in the tree)
    - node: gtsam.Pose3 (the newly sampled node)

    Returns:
    - nearest: gtsam.Pose3 (the node in the tree which is CLOSEST to the newly sampled node)
    - index: int (the index of the closest node, so we can keep track of the parent)
    ...

nearest = None
index = None

```

```
##### Student code here #####
min_distance = float('inf')

# Looping through all nodes and find the closest one
for i, pose in enumerate(rrt):
    dist = helpers.distance_between_poses(pose, node)
    if dist < min_distance:
        min_distance = dist
        nearest = pose
        index = i

#raise NotImplementedError("find_nearest_pose is not implemented")

##### End student code #####
return nearest, index
```

```
suite = unittest.TestSuite()
suite.addTest(TestSteeringWithTerminalVelocity('test_find_nearest_pose'))

unittest.TextTestRunner().run(suite)

.
-----
Ran 1 test in 0.007s

OK
<unittest.runner.TextTestResult run=1 errors=0 failures=0>
```

Finally, we can rewrite the `steer()` function used in the RRT algorithm. We want to fly the drone with terminal velocity in the direction of the new random node for a specific duration, and return the node we reach.

In this part, we can assume that we can apply the thrust in the direction we want to steer in, and we always apply the maximum thrust of $20N$. You may want to compute the following:

- direction of travel
- new attitude of the drone using `get_new_attitude()` method from `helpers`
- force vector using `compute_force()`
- terminal velocity using `compute_terminal_velocity()`
- new position by applying the velocity for the given duration

```
help(helpers.get_new_attitude)

Help on function get_new_attitude in module helpers_obstacles:

get_new_attitude(current: gtsam.gtsam.Pose3, direction: <function Point3 at 0x7a88081dccc0>)
```

```
# TODO 11
def steer_with_terminal_velocity(current: gtsam.Pose3, target: gtsam.Pose3, duration: float = 0.1) -> gtsam.Pose3:
    ...

We need to find a short steering from the current pose toward the target pose.
We fly the drone for a small duration in the direction of the target node at
the terminal velocity with maximum possible thrust of 20N.

Arguments:
- current: gtsam.Pose3, the current pose of the drone
- target: gtsam.Pose3, the target pose of the drone
- duration: float, the duration of the flight, default: 0.1

Returns:
- steer_node: gtsam.Pose3, the new node reached by the drone
...

steer_node = None

##### Student code here #####
# Calculating direction we want to go
current_pos = np.array(current.translation())
target_pos = np.array(target.translation())

direction = target_pos - current_pos
direction_norm = np.linalg.norm(direction)
```

```

# Handling case where we're already at target
if direction_norm < 1e-6:
    return current

# Normalize the direction
direction_unit = direction / direction_norm

# Getting a new attitude pointing in the travel direction
# Passig the current pose and convert direction_unit to gtsam.Point3
new_attitude = helpers.get_new_attitude(current, gtsam.Point3(direction_unit[0], direction_unit[1], direction_unit[2]))

# Calculating force with maximum thrust
max_thrust = 20.0
force = compute_force(new_attitude, max_thrust)

# Getting the terminal velocity
terminal_vel = compute_terminal_velocity(force)

# Calculating new position after flying for 'duration' seconds
velocity_array = np.array([terminal_vel[0], terminal_vel[1], terminal_vel[2]])
new_position = current_pos + velocity_array * duration

# Creating the steer node
steer_node = gtsam.Pose3(new_attitude, gtsam.Point3(new_position[0], new_position[1], new_position[2]))

# raise NotImplementedError("steer_with_terminal_velocity is not implemented")

##### End student code #####
return steer_node

```

```
help(steer_with_terminal_velocity)
```

Help on function steer_with_terminal_velocity in module __main__:

```
steer_with_terminal_velocity(current: gtsam.gtsam.Pose3, target: gtsam.gtsam.Pose3, duration: float = 0.1) -> gtsam.gtsam.Pose3
    We need to find a short steering from the current pose toward the target pose.
    We fly the drone for a small duration in the direction of the target node at
    the terminal velocity with maximum possible thrust of 20N.
```

Arguments:

- current: gtsam.Pose3, the current pose of the drone
- target: gtsam.Pose3, the target pose of the drone
- duration: float, the duration of the flight, default: 0.1

Returns:

- steer_node: gtsam.Pose3, the new node reached by the drone

```
suite = unittest.TestSuite()
suite.addTest(TestSteeringWithTerminalVelocity('test_steer_with_terminal_velocity'))
```

```
unittest.TextTestRunner().run(suite)
```

```
.
-----
Ran 1 test in 0.008s
```

```
OK
<unittest.runner.TextTestResult run=1 errors=0 failures=0>
```

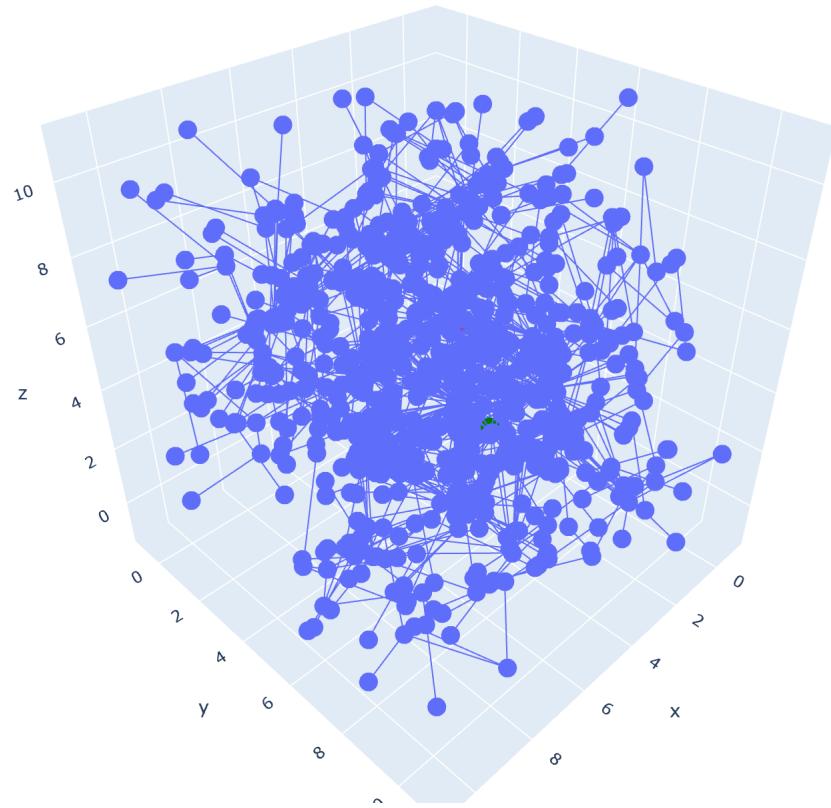
Putting it back into RRT

- Now that we have coded the different components using drone dynamics, let us run the RRT algorithm again with the new steer function!

```
start_rrt_drone = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(1, 2, 3))
target_rrt_drone = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(5, 7, 4))
rrt_drone, parents_rrt_drone = run_rrt(start_rrt_drone, target_rrt_drone, generate_random_pose, steer_with_terminal_velocity,
                                         helpers.distance_between_poses, find_nearest_pose, threshold=1.5)
print("Number of RRT Nodes: ", len(rrt_drone))
```

```
Number of RRT Nodes:  973
```

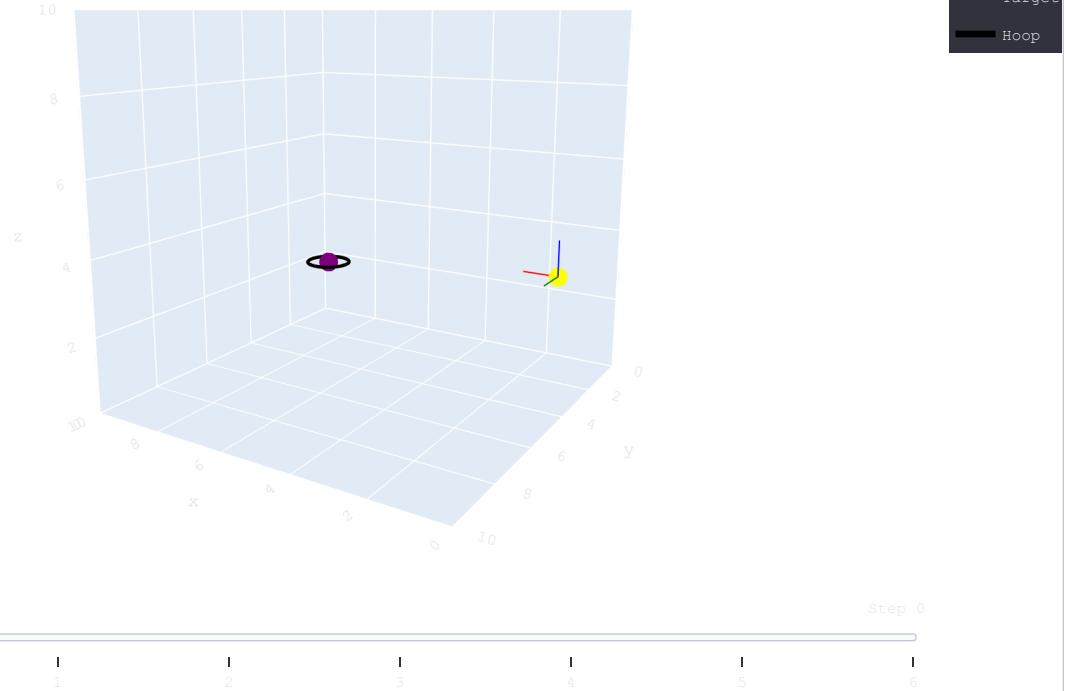
```
helpers.visualize_tree(rrt_drone, parents_rrt_drone, start_rrt_drone, target_rrt_drone)
```



```
path_rrt_drone = get_rrt_path(rrt_drone, parents_rrt_drone)
print("Length of Path: ", len(path_rrt_drone))
print("The path obtained by steering with a terminal velocity:")
helpers.animate_drone_path(path_rrt_drone, start_rrt_drone, target_rrt_drone)
```

Length of Path: 7

The path obtained by steering with a terminal velocity:



Reflection Questions

- Is this `steer` function we coded up realistic in nature? What do you think may be missing, if so?

Response:

No, not really realistic at all. The biggest issue is that we are basically 'teleporting' the orientation of the drone - we instantly point it in whatever direction we want to go. Real drones can't just snap to a new attitude like that; the motors need time to adjust and there is inertia that comes into play. We're also ignoring gravity completely here, which is a huge deal for actual flight. In the calculation of terminal velocity, it was assumed that we obtain maximum speed instantaneously. In reality, there is a time period for acceleration. We don't account for any rotation rate limits - real quadrotors can only yaw/pitch/roll so fast. Lacking wind resistance and other environmental factors. With no consideration of motor response time or control delays. Basically, we are assuming perfect instantaneous control. That does not exist in the real world. Would need to model the actual dynamics with angular velocities, acceleration limits, and a time to change attitude. Also the transition from one orientation to another needs to be smooth, not instant.

▼ Part 4: Steering with more realistic drone dynamics

Two major limitations stand out in our previous approach of steering the drone:

1. We ignored the gravitational force affecting the direction and magnitude of our applied thrust. If you didn't believe this was important, [let Sheldon and Leonard tell you otherwise](#).
2. By assuming we could apply the thrust in the direction we want to steer in, we allowed arbitrary and instantaneous attitude changes at the time of steering.

We are going to address these concerns in this section.

Here's the test suite for this part. Feel free to add your test cases for further testing.

```

class TestRealisticSteer(unittest.TestCase):
    def test_compute_force_with_gravity(self):
        attitude = gtsam.Rot3(
            [0.612372, 0.612372, -0.5],
            [-0.0473672, 0.65974, 0.75],
            [0.789149, -0.435596, 0.433013]
        )
        thrust = 20.0

        expected_force = gtsam.Point3(15.78, -8.71, -1.34)
        actual_force = compute_force_with_gravity(attitude, thrust)

        assert(np.allclose(actual_force, expected_force, atol=1e-2))

    def test_steer(self):
        current_node = gtsam.Pose3(gtsam.Rot3.Yaw(math.radians(90)), gtsam.Point3(1, 2, 3))
        new_node = gtsam.Pose3(gtsam.Rot3.Pitch(math.radians(45)), gtsam.Point3(8, 5, 6))

        expected_steer_node = gtsam.Pose3(gtsam.Rot3(
            [0.17, 0.97, -0.17],
            [-0.96, 0.20, 0.17],
            [0.20, 0.14, 0.97]
        ), gtsam.Point3(1.97, 2.81, 4.49))
        actual_steer_node = steer(current_node, new_node)

        assert(actual_steer_node.equals(expected_steer_node, tol=1e-2))

```

▼ Hovering the drone

To hover, assuming $g = 10 \text{ m/s}^2$, the four rotors have to provide a thrust of $10N$ upwards to compensate for gravity, i.e., $2.5N$ per motor. Of course, we need to be able to accelerate upwards, so let's assume each motor can provide up to double that, i.e., 0 to $5N$. So, while the drone is level, here are some sample accelerations we can deliver (note that the drone has a mass of $1kg$):

- $f_i = 0N$ for $i \in 1..4$: downwards acceleration at -10 m/s^2
- $f_i = 2.5N$ for $i \in 1..4$: stable hover 0 m/s^2
- $f_i = 5N$ for $i \in 1..4$: upwards acceleration at 10 m/s^2

▼ Correcting force vector

We will now correct our implementation of `compute_force` to incorporate the effect of the gravitational force. We can assume $g = 10 \text{ m/s}^2$.

```

# TODO 12
def compute_force_with_gravity(attitude: gtsam.Rot3, thrust: float, mass: float = 1.0) -> gtsam.Point3:
    """
    Computes the net force vector given attitude and thrust in the body frame
    by adjusting for the downwards weight force.

    Arguments:
    - attitude: gtsam.Rot3, nRb for the drone
    - thrust: float, the upwards thrust produced by the 4 rotors
    - mass: float, the mass of the drone, default: 1.0

    Returns:
    - force: gtsam.Point3, the resultant force vector
    """

    force = None
    g = 10.0 # m/s^2

    ##### Student code here #####
    # Thrust force in navigation frame (from body frame)
    thrust_body = np.array([0, 0, thrust])
    thrust_nav = attitude.matrix() @ thrust_body

    # Gravity force (always pointing down in navigation frame)
    gravity_force = np.array([0, 0, -mass * g])

    # Net force is thrust minus gravity

```

```

net_force = thrust_nav + gravity_force

force = gtsam.Point3(net_force[0], net_force[1], net_force[2])

#raise NotImplementedError("compute_force_with_gravity is not implemented")

##### End student code #####
return force

```

```

suite = unittest.TestSuite()
suite.addTest(TestRealisticSteer('test_compute_force_with_gravity'))

unittest.TextTestRunner().run(suite)

.
-----
Ran 1 test in 0.002s

OK
<unittest.runner.TextTestResult run=1 errors=0 failures=0>

```

▼ Reflection questions

Since we now have the correct drone dynamics in place, let's answer a few reflection questions.

Play with the configurations of yaw, pitch, roll and thrust (in the cell below) and answer these questions for each configuration:

1. In which direction is the drone flying in the navigation frame?
2. Is the drone flying upwards, downwards or maintaining level flight?
3. In what direction is the thrust applied in the navigation frame?
4. What is the speed of the drone?
5. What direction is the front of the drone facing in the navigation frame?

Feel free to add your own code to the cell below, if you need it.

Configurations: All angles are in degrees and thrust is in Newtons.

Yaw	Pitch	Roll	Thrust
90	0	0	10
45	45	45	20
30	0	45	15
10	30	30	0

Double-click (or enter) to edit

▼ Response:

Configuration 1: Yaw=90°, Pitch=0°, Roll=0°, Thrust=10N

1. Direction of flight: The drone is hovering in place, not really moving horizontally (maybe slight drift)
2. Flight mode: Level flight - thrust perfectly balances gravity
3. Thrust direction: Pointing straight up in the navigation frame
4. Speed: Pretty much zero or very slow
5. Front facing: East direction (90° yaw means rotated from north to east)

Configuration 2: Yaw=45°, Pitch=45°, Roll=45°, Thrust=20N

1. Direction of flight: This is complicated - moving northeast and upward with some lateral component from the roll
2. Flight mode: Flying upwards - thrust is way more than needed for hover
3. Thrust direction: Angled toward northeast and up
4. Speed: Pretty fast, probably around 15-20 m/s since we're at max thrust
5. Front facing: Northeast direction (45° yaw)

Configuration 3: Yaw=30°, Pitch=0°, Roll=45°, Thrust=15N

1. Direction of flight: Moving sideways to the left (because of the roll) and slightly up
2. Flight mode: Climbing slightly - thrust is above hover level

3. Thrust direction: Mostly up but tilted to the left due to roll
4. Speed: Moderate, maybe 8-10 m/s
5. Front facing: 30° east of north

Configuration 4: Yaw=10°, Pitch=30°, Roll=30°, Thrust=ON

1. Direction of flight: Falling down - no thrust to fight gravity
2. Flight mode: Definitely going downwards, basically in freefall
3. Thrust direction: No thrust applied, only gravity pulling down
4. Speed: Accelerating downward at 10 m/s² until terminal velocity from drag
5. Front facing: Slightly east of north (10° yaw), but tilted forward and to the side

General observations:

The main thing I noticed is that pitch controls forward/backward movement, roll controls left/right movement, and yaw just spins the drone around. When thrust equals 10N we hover, less than that we fall, more than that we climb. RetryTo run code, enable code execution and file creation in Settings > Capabilities.

```
# Use for the reflection questions above

# Controllers
yaw = math.radians(45)
pitch = math.radians(45)
roll = math.radians(45)
thrust = 20

# Computing drone velocity
nRb = compute_attitude_from_ypr(yaw, pitch, roll)
net_force = compute_force_with_gravity(nRb, thrust)
terminal_velocity = compute_terminal_velocity(net_force)

print(f"Orientation in navigation frame (nRb):\n {np.round(nRb.matrix(), 2)}")
print(f"Net force:\n {np.round(net_force, 2)}")
print(f"Terminal velocity:\n {np.round(terminal_velocity, 2)}")

Orientation in navigation frame (nRb):
[[ 0.5 -0.15  0.85]
 [ 0.5   0.85 -0.15]
 [-0.71  0.5   0.5 ]]
Net force:
[17.07 -2.93  0. ]
Terminal velocity:
[20.04 -8.3   0. ]
```

▼ Correcting drone attitude while steering

In this section, we will implement a more realistic version of the `steer()` function. We cannot just rotate the drone into any direction of travel. We limit the yaw, pitch and roll rotations of the drone to $[-10^\circ, 10^\circ]$ for a more realistic instantaneous change in attitude. We also attempt to find a thrust value that would take us closest to the new random node, for each yaw, pitch and roll.

To make it more simple, we select a yaw from 3 values: $[-10^\circ, 0^\circ, 10^\circ]$, a pitch from 3 values: $[-10^\circ, 0^\circ, 10^\circ]$, a roll from 3 values: $[-10^\circ, 0^\circ, 10^\circ]$ and a thrust from 4 values: $[5, 10, 15, 20]$. For each of the 108 combinations, we find the node we can steer to, by finding the new attitude and applying the terminal velocity computed using the updated force function for a specific duration. We return the node that is closest to the target node we are steering towards.

▼ Hint for calculating the new attitude:

The yaw, pitch and roll rotations of $[-10^\circ, 0^\circ, 10^\circ]$ are in the body frame. You can use the existing attitude and this rotation in body frame to compute the new attitude.

$$R_{b_1}^n = R_{b_0}^n R_{b_1}^{b_0}$$

```
# TODO 13
def steer(current: gtsam.Pose3, target: gtsam.Pose3, duration = 0.1):
    ...
    Steering with limits on rotation change and thrust values
    using force computations with gravity
    ...
```

```

steer_node = None
yaw_values = [-10, 0, 10]
pitch_values = [-10, 0, 10]
roll_values = [-10, 0, 10]
thrust_values = [5, 10, 15, 20]

##### Student code here #####
best_node = None
min_distance = float('inf')

current_pos = np.array(current.translation())
current_rot = current.rotation()

# Try all combinations of controls
for dyaw_deg in yaw_values:
    for dpitch_deg in pitch_values:
        for droll_deg in roll_values:
            for thrust in thrust_values:
                # Convert to radians
                dyaw = dyaw_deg * np.pi / 180
                dpitch = dpitch_deg * np.pi / 180
                droll = droll_deg * np.pi / 180

                # Creating rotation change in body frame
                delta_rotation = compute_attitude_from_ypr(dyaw, dpitch, droll)

                # Applying to current attitude: new = current * delta
                new_rot = current_rot.compose(delta_rotation)

                # Calculating force with gravity
                force = compute_force_with_gravity(new_rot, thrust)

                # Getting terminal velocity
                vel = compute_terminal_velocity(force)

                # Calculating new position after flying for 'duration' seconds
                vel_array = np.array([vel[0], vel[1], vel[2]])
                new_pos = current_pos + vel_array * duration

                # Creating candidate node
                candidate = gtsam.Pose3(new_rot, gtsam.Point3(new_pos[0], new_pos[1], new_pos[2]))

                # Checking if it is closer to target
                dist = helpers.distance_between_poses(candidate, target)
                if dist < min_distance:
                    min_distance = dist
                    best_node = candidate

steer_node = best_node

##### End student code #####
return steer_node

```

```

suite = unittest.TestSuite()
suite.addTest(TestRealisticSteer('test_steer'))

unittest.TextTestRunner().run(suite)

.
-----
Ran 1 test in 0.016s

OK
<unittest.runner.TextTestResult run=1 errors=0 failures=0>

```

▼ RRT with drone dynamics!

Now that we have coded the appropriate drone dynamics, let us plug the `steer()` method back into our RRT function

```

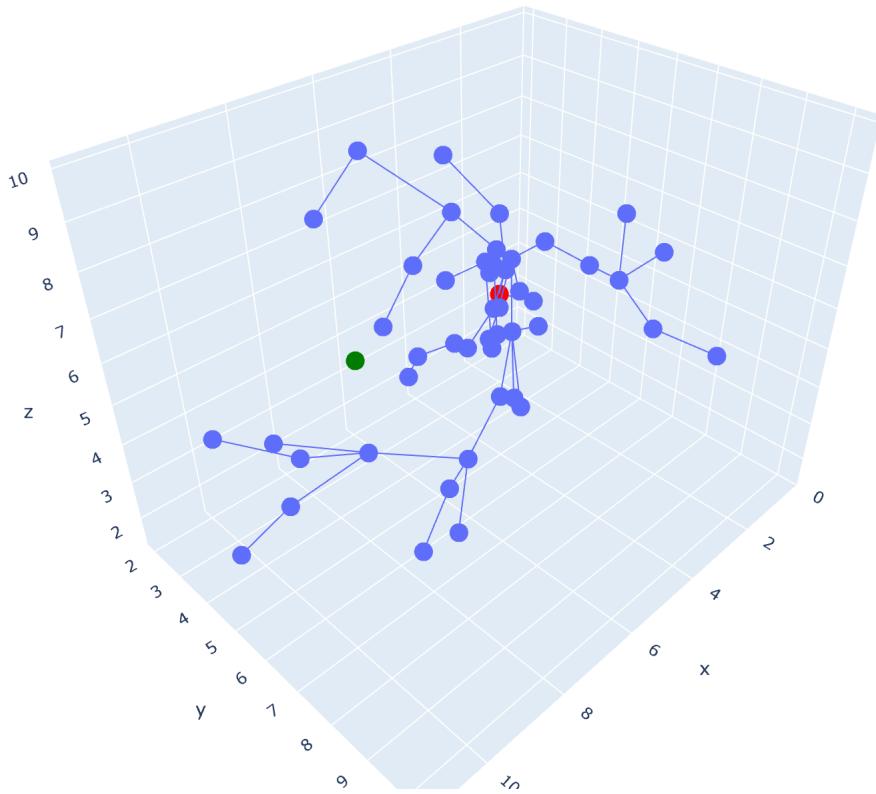
start_rrt_drone_realistic = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(1, 2, 3))
target_rrt_drone_realistic = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(8, 5, 6))
rrt_drone_realistic, parents_rrt_drone_realistic = run_rrt(start_rrt_drone_realistic, target_rrt_drone_realistic,
    generate_random_pose, steer, helpers.distance_between_poses,

```

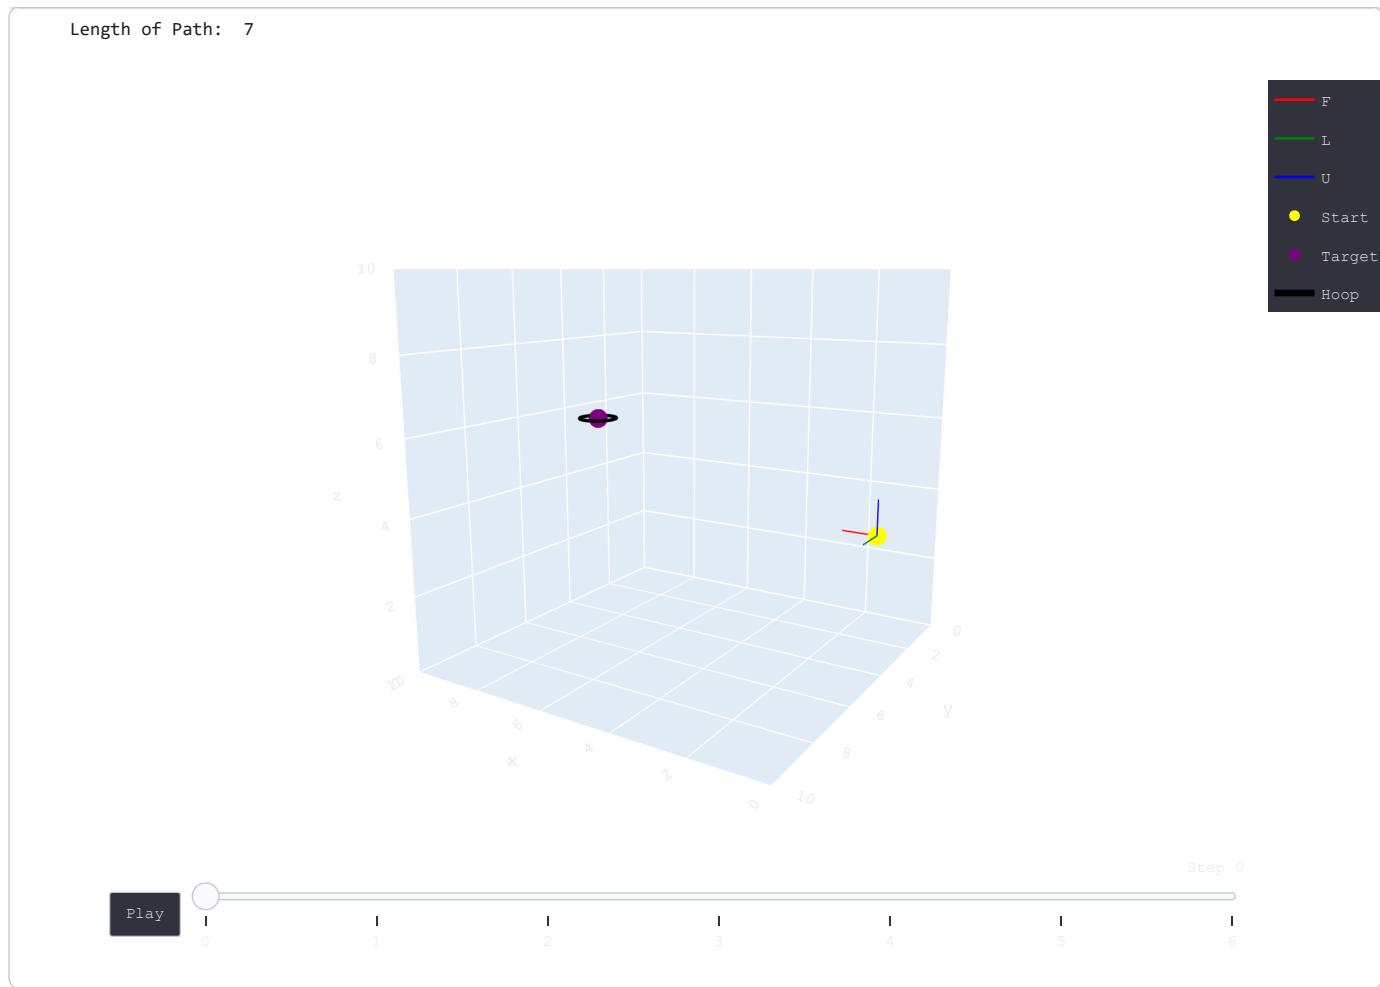
```
find_nearest_pose, threshold=1.5)  
print("Nodes in RRT Tree: ", len(rrt_drone_realistic))  
Nodes in RRT Tree: 51
```

Let us see how our new RRT algorithm now performs with the same start and target poses.

```
helpers.visualize_tree(rrt_drone_realistic, parents_rrt_drone_realistic, start_rrt_drone_realistic, target_rrt_drone_realistic)
```



```
path_rrt_drone_realistic = get_rrt_path(rrt_drone_realistic, parents_rrt_drone_realistic)  
print("Length of Path: ", len(path_rrt_drone_realistic))  
helpers.animate_drone_path(path_rrt_drone_realistic, start_rrt_drone_realistic, target_rrt_drone_realistic)
```



Reflection Questions

- Do you think the drone can fly faster than our current implementation using these realistic dynamics? If yes, how?

Response:

Yes, it could definitely go faster with some changes. Right now we're always trying to minimize changes in thrust and orientation, which makes everything smooth but slow.

- We could be more aggressive with the pitch angles; leaning forward further would yield higher horizontal speed
- Duration parameter is set to 0.1s which is pretty conservative, could use longer prediction horizons probably
- We limit ourselves to ± 10 degrees per step for yaw/pitch/roll; this is realistic yet conservative, as racing drones can take more aggressive moves
- Thrust values could be optimized better - as of now we simply try 5, 10, 15, 20 but perhaps intermediate values may help
- The steer function tries all 108 combinations which is slow. A smarter heuristic could pick better options faster
- We're also not really planning ahead - just greedily picking the best immediate move toward target
- In real drone racing, they pre-plan optimal trajectories considering the full path, not just the next step
- Could use higher thrust values when needed though this is limited by the motor specifications
- In general, we're flying "safely" when we could be flying "fast" if we relaxed some constraints

▼ Part 5: Drone Racing in Free environment!

- Alright folks, now that we've successfully implemented the RRT algorithm whilst incorporating drone dynamics, it is now time to take your drones to the **NEU 5335 EASY Drone Racing Challenge!!** ARE YOU READY?!?! Check out this exhilarating racing video to get your adrenaline rushing!



- The idea is simple; there are 4 hoops in the NEU Racing Track, which are given to you as `targets`. Your objective is to fly through those hoops in the **same order** and reach the final treasure, which is present at the center of the last hoop in the list.
- We know that RRT plans a path from one start point to one destination point. Since we have multiple hoops to fly through, we can consider them as **intermediate goal points**.
- So you can first use the RRT algorithm to plan a path from the start location to the first hoop (an intermediate goal), then use RRT again to plan your path from that point to the next hoop, and so on and so forth!
- As you would've noticed in Part 4, the path generated by RRT gets us really close to the goal, but doesn't get us exactly there. So we've provided you with a magical function `pass_through_the_hoop`, which performs some wizardry and gets your drone through the hoop 😊
- Now let's have some fun racing the drone for the easy environment and later we will build it for a harder one!

You need to perform RRT with a set of intermediate goal points, use `pass_through_the_hoop` to successfully traverse through the hoops, and obtain your final treasure!

When you call RRT with `run_rrt`, make sure you pass the correct functions as arguments! **You can play around with the threshold, but do NOT exceed a threshold value of 3**

```
# TODO 14
def drone_racing_rrt(start: gtsam.Pose3, targets: List[gtsam.Pose3]) -> List[gtsam.Pose3]:
    ...
    Runs RRT multiple times from start to each target in order
    Note: Since the drone can only reach near the target node, we add
    multiple points at the end of each sub-path to pass through the hoop.
    Use `helpers.pass_through_the_hoop(target, path)` after calculating an RRT path
    to each target for appending the relevant points.

    Arguments:
        - start: gtsam.Pose3, initial position of the drone
        - targets: List[gtsam.Pose3], RRT targets

    Returns:
        - drone_path: List[gtsam.Pose3], entire path from start to last hoop
        ...

    drone_path = []

##### Student code here #####
# Keep track of current position
current_start = start

# Processing each hoop target
for i, target in enumerate(targets):
    # Run RRT to reach this hoop
    rrt, parents = run_rrt(
        current_start,
        target,
        generate_random_pose,
        steer_with_terminal_velocity, # Use simpler steer
        helpers.distance_between_poses,
        find_nearest_pose,
        threshold=2.0
    )
    ...
    # Append the points from the RRT path to the overall drone path
    drone_path += rrt[1:-1] + [target]

# The final path starts at the initial start and ends at the last target
drone_path = [start] + drone_path + [targets[-1]]
```

```
# Extracting path from tree
path = get_rrt_path(rrt, parents)

# Checking if we got a valid path
if path is None or len(path) == 0:
    print(f"Failed to find path to hoop {i}")
    continue

# Trying to pass through the hoop
path_with_hoop = helpers.pass_through_the_hoop(target, path)

# Checking if pass_through_the_hoop returned a valid result
if path_with_hoop is None or len(path_with_hoop) == 0:
    # Fallback: just use the RRT path without hoop waypoints
    path_with_hoop = path

# Adding to overall drone path
if len(drone_path) > 0:
    # Skipping the first element to avoid duplicate
    drone_path.extend(path_with_hoop[1:])
else:
    # First segment, add everything
    drone_path.extend(path_with_hoop)

# Updating starting position for next iteration
current_start = path_with_hoop[-1]

##### End student code #####
return drone_path
```

```
start_rrt_drone_race = gtsam.Pose3(r=gtsam.Rot3(), t=gtsam.Point3(1, 3, 8))
targets_rrt_drone_race = helpers.get_targets()
path_rrt_drone_race = drone_racing_rrt(start_rrt_drone_race, targets_rrt_drone_race)
```

```
helpers.drone_racing_path(helpers.get_hoops(), start_rrt_drone_race, path_rrt_drone_race)
print("Length of Drone Racing Path: ", len(path_rrt_drone_race))
```

Part 6.1: RRT with Obstacle Avoidance



Why Obstacles Matter for Trajectory Optimization

So far, we've successfully navigated the drone through hoops in an **obstacle-free environment**. While our RRT paths work, they have major limitations:

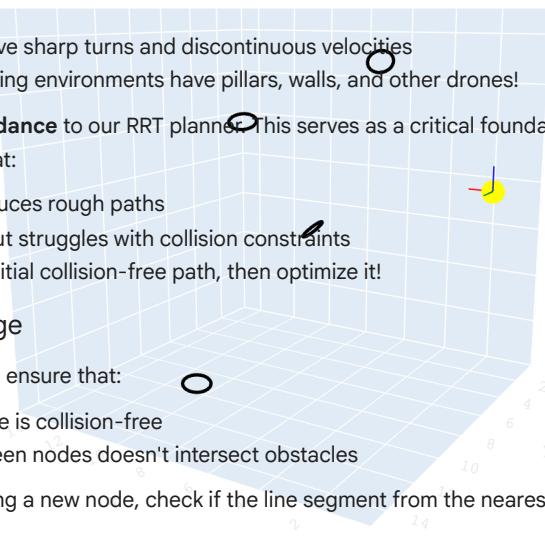
- 1. Jerky trajectories:** RRT paths have sharp turns and discontinuous velocities.
- 2. No obstacle avoidance:** Real racing environments have pillars, walls, and other drones!

In this section, we'll add **obstacle avoidance** to our RRT planner. This serves as a critical foundation for **Part 7 (Trajectory Optimization)**, where we'll discover that:

- RRT can avoid obstacles but produces rough paths
- Optimization can smooth paths but struggles with collision constraints
- The best solution:** Use RRT for initial collision-free path, then optimize it!

The Collision Checking Challenge

When planning with obstacles, we must ensure that:



Key modification to RRT: Before adding a new node, check if the line segment from the nearest node to the new node is collision-free.

Obstacle Types

We'll work with one obstacle primitive (defined in `helpers.py`) which is a `SphereObstacle(center, radius, name)`: Circular obstacles (pillars, balloons)

The `helpers` module provides collision checking utilities:

- `check_segment_collision(p1, p2, obstacles)`: Returns `True` if segment collides
- `check_path_collision(path, obstacles)`: Validates entire path

```
# TODO 15: RRT with Obstacle Avoidance
def run_rrt_with_obstacles(start, target, generate_random_node, steer, distance,
                           find_nearest_node, threshold, obstacles: List = None):
    ...
    Run RRT with collision checking to avoid obstacles.

    This is nearly identical to your `run_rrt` from TODO 5, with ONE critical addition:
    Before adding a new node to the tree, verify that the path segment from the nearest
    node to the new node is collision-free.

    Hints:
    - Start by using your `run_rrt` implementation from TODO 5
    - Use `helpers.check_segment_collision(p_nearest, p_new, obstacles)`
    - If collision detected, use `continue` to skip to next iteration
    - Increase max_iterations to 5000 (obstacles make planning harder)

    **Testing:**
    - With no obstacles: should behave like regular RRT
    - With obstacles: should route around them, may need more iterations

    Arguments:
    - start: gtsam.Point3 or gtsam.Pose3, starting configuration
    - target: gtsam.Point3 or gtsam.Pose3, goal configuration
    - generate_random_node: function to sample random configurations
    - steer: function to steer toward target
    - distance: function to compute distance between configurations
    - find_nearest_node: function to find closest node in tree
    - threshold: float, distance threshold for goal reaching
    - obstacles: List[helpers.Obstacle], obstacles to avoid (default: None/empty)

    Returns:
    - rrt: List, the RRT tree of configurations
    - parents: List[int], parent indices for path reconstruction
    ...
```

```

rrt = None
parents = None

##### Student code here #####
rrt = []
parents = []
max_iterations = 5000 # More iterations for obstacle environments

rrt.append(start)
parents.append(-1)

# Handle case with no obstacles
if obstacles is None:
    obstacles = []

for i in range(max_iterations):
    # Step 1: Sampling random node with goal bias
    random_node = generate_random_node(target)

    # Step 2: Finding nearest node in tree
    nearest_node, nearest_idx = find_nearest_node(rrt, random_node)

    # Step 3: Steering toward the random node
    new_node = steer(nearest_node, random_node)

    # Step 4: CHECKING FOR COLLISIONS (new part!)
    # Extracting the positions for collision checking
    if hasattr(nearest_node, 'translation'):
        # It's a Pose3
        p_nearest = np.array(nearest_node.translation())
        p_new = np.array(new_node.translation())
    else:
        # It's a Point3
        p_nearest = np.array([nearest_node[0], nearest_node[1], nearest_node[2]])
        p_new = np.array([new_node[0], new_node[1], new_node[2]])

    # Checking if line segment collides with any obstacle
    has_collision = helpers.check_segment_collision(p_nearest, p_new, obstacles)

    # If collision is detected, skip this node and try another
    if has_collision:
        continue

    # Step 5: Safe to add this node to the tree
    rrt.append(new_node)
    parents.append(nearest_idx)

    # Step 6: Checking if we reached the goal
    dist_to_goal = distance(new_node, target)
    if dist_to_goal < threshold:
        # Success!
        return rrt, parents

# Return tree even if we didn't reach goal (best effort)

#####
# End student code #####
return rrt, parents

```

Let's test our RRT with obstacles implementation on a simple scenario before tackling the full racing circuit.

```

# Test RRT with obstacles on simple scenario
print("Testing RRT with obstacle avoidance...")

# Simple test: navigate around a single sphere obstacle
start_obs = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(2, 2, 5))
target_obs = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(8, 8, 8))

# Create obstacle in the middle
obstacles_simple = [
    helpers.SphereObstacle(center=[5, 5, 6.5], radius=1.5, name="Central Pillar")
]

# Run RRT with obstacles

```

```
rrt_obs, parents_obs = run_rrt_with_obstacles(
    start=start_obs,
    target=target_obs,
    generate_random_node=generate_random_pose,
    steer=steer,
    distance=helpers.distance_between_poses,
    find_nearest_node=find_nearest_pose,
    threshold=1.5,
    obstacles=obstacles_simple
)

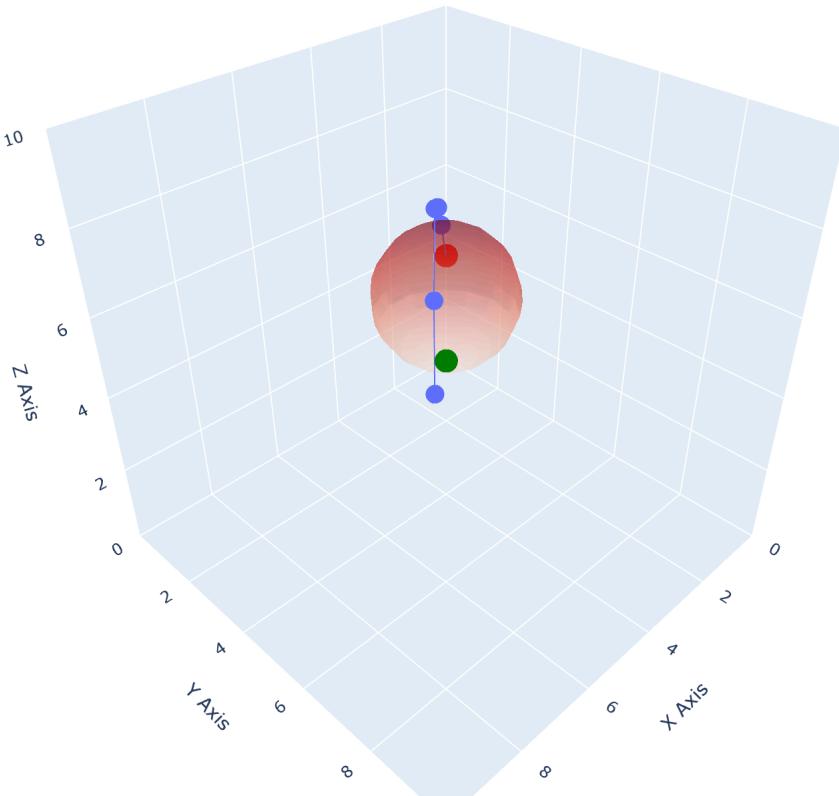
print(f"RRT with obstacles completed: {len(rrt_obs)} nodes")

# Extract and verify path
path_obs = get_rrt_path(rrt_obs, parents_obs)
print(f"Path length: {len(path_obs)} waypoints")

# Verify no collisions
has_collision, _ = helpers.check_path_collision(path_obs, obstacles_simple)
if not has_collision:
    print(f"RRT is Collision-free!")
else:
    print(f"RRT may have collisions")
```

```
Testing RRT with obstacle avoidance...
RRT with obstacles completed: 12 nodes
Path length: 6 waypoints
RRT is Collision-free!
```

```
# Visualize RRT with obstacles
helpers.visualize_path_with_obstacles(
    path_obs, start_obs, target_obs, obstacles_simple
)
```



❖ PART 6.2: Drone Racing with Obstacles — The Real Challenge

Alright, this is where things get *real*. So far, you've been navigating clean airspace - no walls, no distractions. But now? We're dropping you into a proper racing arena, full of hoops **and** obstacles. Think of it as "Formula 1 for drones"... but with sphere obstacles termed as chaos 😅.

Let's see if your planner can handle it.

The Scenario

Here's what we're setting up:

- **4 racing hoops** placed in the environment (same as Part 5).
- **Random obstacles** scattered throughout the course.
- **Goal:** Your drone must pass through **every hoop** in order, without colliding with anything.

The Setup

We'll feed this into our RRT planner:

- Each hoop acts as a **waypoint goal** — the RRT will plan from one hoop to the next.
- The environment now has **obstacle volumes** that the planner must avoid.
- The RRT node expansion logic is the same, but now every new edge must pass the *collision check* before being added.

Why This Is Tough (and Fun)

Let's break down what makes this scenario interesting (and frustrating in the best way):

- **Jagged RRT paths:** With obstacles everywhere, RRT has to wiggle around them. The result? Paths that look like spaghetti, not the clean racing line you want.

- **More nodes = slower planning:** The planner will need more samples to find valid routes between hoops. Be patient - this is the trade-off between complexity and realism.
- **Optimization becomes essential:** This is *the reason* we care about trajectory optimization (discussed in Part 7). RRT can get you there, but only roughly. To make it *flyable* for a real drone, we'll soon smooth and optimize it into a continuous, dynamically feasible trajectory.

Pro tip: Rewatch that racing video from earlier - the adrenaline will make debugging much less painful.

```
# TODO 16: Drone Racing with Obstacle Avoidance
def drone_racing_rrt_with_obstacles(start: gtsam.Pose3, targets: List[gtsam.Pose3],
                                      obstacles: List = None) -> List[gtsam.Pose3]:
    ...
    Navigate through racing hoops while avoiding obstacles.
    ...

    drone_path = []

    ##### Student code here #####
    # Handle case with no obstacles
    if obstacles is None:
        obstacles = []

    # Start from initial position
    current_start = start

    # Processing each hoop target
    for i, target in enumerate(targets):
        print(f"Planning to hoop {i}...")

        path_found = False
        best_path = None
        best_distance = float('inf')

        # Trying multiple attempts with varying parameters
        for attempt in range(10):
            # Gradually relax threshold
            current_threshold = 2.0 + (attempt * 0.5)

            # Run RRT with obstacle avoidance
            rrt, parents = run_rrt_with_obstacles(
                current_start,
                target,
                generate_random_pose,
                steer,
                helpers.distance_between_poses,
                find_nearest_pose,
                threshold=current_threshold,
                obstacles=obstacles
            )

            # Extracting path
            path = get_rrt_path(rrt, parents)

            # Checking if we got a valid path
            if path is None or len(path) == 0:
                continue

            # Checking how close we got
            final_dist = helpers.distance_between_poses(path[-1], target)

            # Keeping track of the best attempt
            if final_dist < best_distance:
                best_distance = final_dist
                best_path = path

            # Trying to pass through hoop
            try:
                processed_path = helpers.pass_through_the_hoop(target, path)

                if processed_path is not None and len(processed_path) > 0:
                    # Success!
                    path = processed_path
            except Exception as e:
                print(f"Exception occurred: {e}")

    return path
```

```

        path_found = True
        print(f" Attempt {attempt+1}: Success!")
        break
    except:
        # If pass_through_hoop fails, continue trying
        continue

    # If pass_through_hoop never worked, use best path we found
    if not path_found:
        if best_path is not None and len(best_path) > 0:
            print(f" Using best path (distance {best_distance:.2f} from hoop)")
            path = best_path
            # Manually add a point at the target to simulate passing through
            path.append(target)
        else:
            print(f"Failed to reach hoop {i}")
            break

    # Adding this segment to overall path
    if len(drone_path) > 0:
        drone_path.extend(path[1:])
    else:
        drone_path.extend(path)

    # Updating starting position
    current_start = path[-1]

#####
# End student code #####
return drone_path

```

Now let's run the full drone racing with obstacles!

```

# Run drone racing with obstacles!
print("Starting drone racing with obstacles...")

start_racing_obs = gtsam.Pose3(r=gtsam.Rot3(), t=gtsam.Point3(1, 3, 8))
targets_racing_obs = helpers.get_targets()

# Get obstacle configuration (easy difficulty)
obstacles_racing = helpers.get_obstacles_easy()
print(f"Racing with {len(obstacles_racing)} obstacles")

# Execute racing
path_racing_obs = drone_racing_rrt_with_obstacles(
    start_racing_obs,
    targets_racing_obs,
    obstacles_racing
)

Starting drone racing with obstacles...
Racing with 2 obstacles
Planning to hoop 0...
    Using best path (distance 1.52 from hoop)
Planning to hoop 1...
    Using best path (distance 1.71 from hoop)
Planning to hoop 2...
    Using best path (distance 1.57 from hoop)
Planning to hoop 3...
    Using best path (distance 1.84 from hoop)

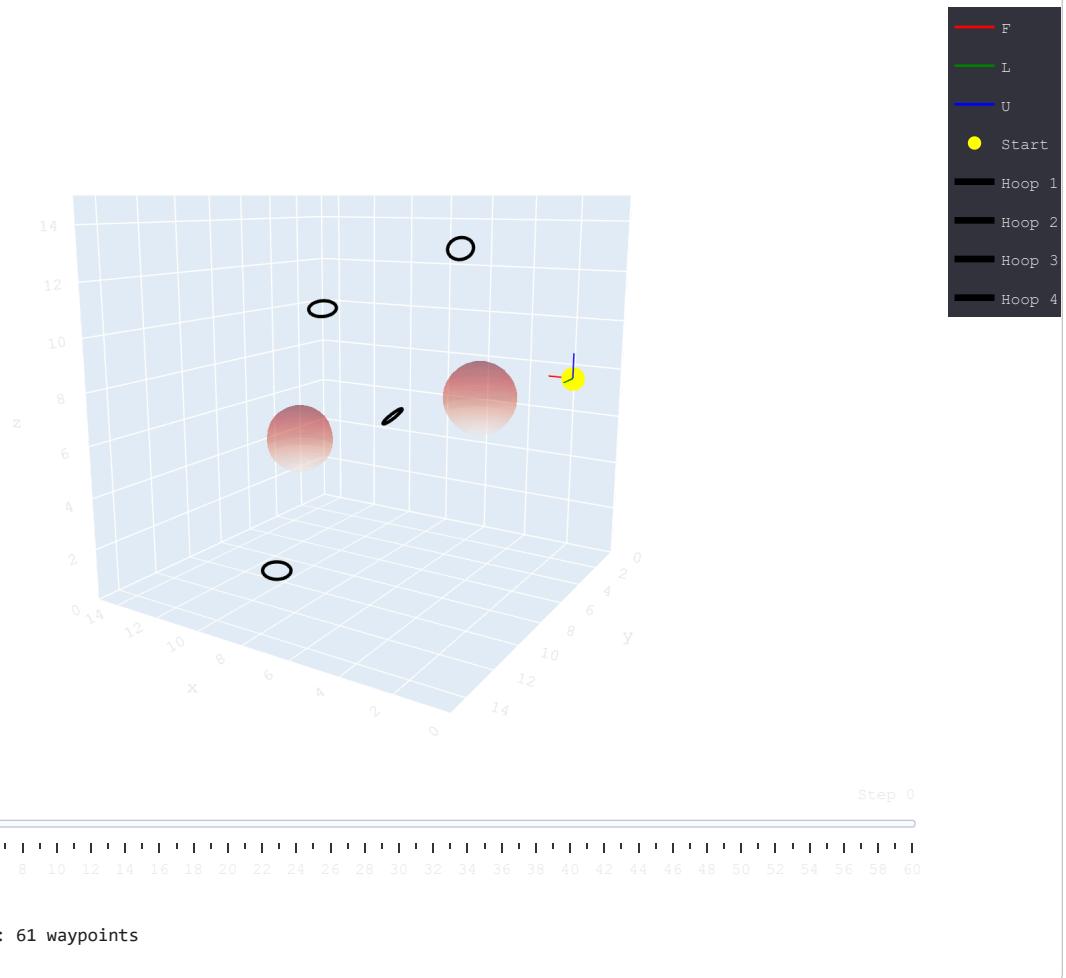
# Visualize racing path with obstacles
helpers.drone_racing_path_with_obstacles(
    hoops=helpers.get_hoops(),
    start=start_racing_obs,
    path=path_racing_obs,
    obstacles=obstacles_racing
)

print(f"\nTotal racing path length: {len(path_racing_obs)} waypoints")

# Verify collision-free
has_collision, _ = helpers.check_path_collision(path_racing_obs, obstacles_racing)
if not has_collision:
    print("RRT is Collision-free!")

```

```
else:
    print(f"RRT may have collisions")
```



Reflection Questions

- What do you think can be some ways to optimize the trajectory taken by the drone? Keep in mind that instantaneous rotation of the drone is limited to -10 to 10 degrees in yaw, pitch and roll.

Response:

- RRT path looks rather messy, with far too many random turns and zigzags.
- We should smooth things out, perhaps fit some curves through the points rather than sharp corners everywhere. All the jerky movements are probably terrible for battery life and could damage the motors over time. It makes more sense to look at the whole path and optimize it entirely, rather than just piece together segments.
- There's definitely redundant waypoints we could cut out - some nodes don't really contribute much. Instead of finding any path, we should try to minimize flight time. Since rotation is limited to ± 10 degree per step, we want to smooth out attitude changes over time so that they are gradual.
- Could throw an optimization algorithm at it to tweak the waypoints and improve overall performance
- Staying close to 10N thrust when possible would help with efficiency
- Speed is all over the place right now; sometimes fast, sometimes slow. Should try to keep it more consistent. Probably could take tighter turns around obstacles without actually hitting them. Something like trajectory optimization would help penalize anything wasteful, like too much thrust or crazy spinning. Just need to ensure that we enforce the physics constraints and the ± 10 degree limits so it stays realistic.

▼ We Need Trajectory Optimization!!

Look at your racing path visualization. You should notice several issues:

Problems with RRT Paths:

1. **Jagged trajectories:** Sharp turns and discontinuous velocities
2. **Inefficient routing:** Takes longer routes to avoid obstacles conservatively
3. **Non-smooth controls:** Would cause jerky flight on a real drone
4. **Slow planning:** Thousands of nodes needed for complex environments
5. **No guarantee of optimality:** Path may be much longer than necessary

This is where trajectory optimization comes in!

In Part 7, we'll learn how to:

- Use RRT to find a **collision-free initial path** (what we just did)
- Feed this path as an **initial guess** to an optimizer
- **Smooth the trajectory** while maintaining collision-free property
- **Minimize energy** and control effort
- **Result:** Smooth racing trajectories!

The key insight: RRT is great at finding feasible paths through complex spaces, but optimization is needed to make them practical for real flight.

Think about it: Would you want to fly on a drone that follows the jagged RRT path, or on one that follows a smooth optimized trajectory?



▼ Part 7: Trajectory Optimization via Direct Transcription

In this final part, we will smooth and optimize the jagged RRT trajectories using **direct transcription** methods. This converts trajectory optimization into a large nonlinear programming (NLP) problem that can be solved efficiently. Read about it here: [Link](#)

Direct Transcription Trajectory Optimization

Overview

RRT produces collision-free paths, but they are often jerky and inefficient. We'll use **direct transcription** to smooth these paths while respecting:

- **Dynamics constraints:** physics must be obeyed between time steps
- **Boundary constraints:** start/goal positions and hoop passage
- **Control limits:** realistic thrust and angular rates
- **Cost minimization:** smooth, efficient trajectories

Mathematical Formulation

State space: $x \in \mathbb{R}^6 = [px, py, pz, \psi, \theta, \phi]$ (position + yaw/pitch/roll)

Control space: $u \in \mathbb{R}^4 = [\Delta\psi, \Delta\theta, \Delta\phi, T]$ (angular changes + thrust)

Decision variables: $z \in \mathbb{R}^{(10N+6)} = [x_0, x_1, \dots, x_N, u_0, u_1, \dots, u_{N-1}]$

Optimization problem:

```

minimize: J(z) = J_thrust + J_angular + J_smoothness + J_gimbal
subject to:
- Dynamics: x_{k+1} = f(x_k, u_k, dt) (terminal velocity model)
- Boundary: x_0 = x_start, x_N at goal, x_h through hoops
- Bounds: |Δψ|, |Δθ|, |Δφ| ≤ 10°, 5 ≤ T ≤ 20
  
```

See [MATHEMATICAL_FORMULATION.md](#) for complete understanding.

```

# Let's make some Test cases for the helper functions we will use to perform Trajectory Optimization!

class TestTrajOptHelpers(unittest.TestCase):
    """Test angle wrapping, packing, and unpacking functions."""
  
```

```

def test_angle_diff_small(self):
    """Test angle_diff with small differences."""
    # Small positive difference
    result = angle_diff(0.1, 0.0)
    self.assertAlmostEqual(result, 0.1, places=6)

    # Small negative difference
    result = angle_diff(0.0, 0.1)
    self.assertAlmostEqual(result, -0.1, places=6)

def test_angle_diff_wrapping(self):
    """Test angle_diff with wrapping around ±π."""
    # Wrapping across +π/-π boundary
    result = angle_diff(np.pi - 0.1, -np.pi + 0.1)
    self.assertAlmostEqual(result, 0.2, places=6)

    # Wrapping the other direction
    result = angle_diff(-np.pi + 0.1, np.pi - 0.1)
    self.assertAlmostEqual(result, -0.2, places=6)

def test_angle_diff_180_degrees(self):
    """Test angle_diff at exactly 180 degrees."""
    # Exactly π apart (ambiguous, but should handle consistently)
    result = angle_diff(0.0, np.pi)
    self.assertTrue(abs(result - np.pi) < 1e-6 or abs(result + np.pi) < 1e-6)

def test_pack_decision_vars(self):
    """Test packing states and controls into flat vector."""
    N = 2
    states = np.array([[1, 2, 3, 0.1, 0.2, 0.3],
                      [4, 5, 6, 0.4, 0.5, 0.6],
                      [7, 8, 9, 0.7, 0.8, 0.9]])  # (N+1) x 6
    controls = np.array([[0.01, 0.02, 0.03, 10],
                        [0.04, 0.05, 0.06, 12]])  # N x 4

    z = pack_decision_vars(states, controls, N)

    # Check size
    self.assertEqual(z.shape[0], 26)  # 10*2 + 6 = 26

    # Check first state
    self.assertTrue(np.allclose(z[:6], [1, 2, 3, 0.1, 0.2, 0.3]))

    # Check last state
    self.assertTrue(np.allclose(z[12:18], [7, 8, 9, 0.7, 0.8, 0.9]))

    # Check first control
    self.assertTrue(np.allclose(z[18:22], [0.01, 0.02, 0.03, 10]))

    # Check second control
    self.assertTrue(np.allclose(z[22:26], [0.04, 0.05, 0.06, 12]))

def test_unpack_decision_vars(self):
    """Test unpacking flat vector into states and controls."""
    N = 2
    # Create a known flat vector
    z = np.array([1, 2, 3, 0.1, 0.2, 0.3,  # state 0
                 4, 5, 6, 0.4, 0.5, 0.6,  # state 1
                 7, 8, 9, 0.7, 0.8, 0.9,  # state 2
                 0.01, 0.02, 0.03, 10,     # control 0
                 0.04, 0.05, 0.06, 12])   # control 1

    states, controls = unpack_decision_vars(z, N)

    # Check shapes
    self.assertEqual(states.shape, (3, 6))
    self.assertEqual(controls.shape, (2, 4))

    # Check first state
    self.assertTrue(np.allclose(states[0, :], [1, 2, 3, 0.1, 0.2, 0.3]))

    # Check last state
    self.assertTrue(np.allclose(states[2, :], [7, 8, 9, 0.7, 0.8, 0.9]))

    # Check controls
    self.assertTrue(np.allclose(controls[0, :], [0.01, 0.02, 0.03, 10]))
    self.assertTrue(np.allclose(controls[1, :], [0.04, 0.05, 0.06, 12]))

```

```
def test_pack_unpack_inverse(self):
    """Test that pack and unpack are inverse operations."""
    N = 3
    # Create random states and controls
    states = np.random.randn(N+1, 6)
    controls = np.random.randn(N, 4)

    # Pack then unpack
    z = pack_decision_vars(states, controls, N)
    states_recovered, controls_recovered = unpack_decision_vars(z, N)

    # Should recover original arrays
    self.assertTrue(np.allclose(states, states_recovered))
    self.assertTrue(np.allclose(controls, controls_recovered))
```

TODO 17: Angle Difference Function
`def angle_diff(angle_to: float, angle_from: float) -> float:`

"""
 Compute shortest angular difference (wrapped to $[-\pi, \pi]$).

1) Why This Matters

Angles wrap around at $\pm\pi$ (180°). The difference between 170° and -170° is NOT 340° , but rather 20° (going the short way arc). This function is CRITICAL for all angle-related costs and constraints in optimization.

2) Mathematical Background

For angles $\alpha, \beta \in [-\pi, \pi]$:

$\Delta = \text{wrap}(\alpha - \beta)$ where $\text{wrap}(\theta)$ maps θ to $[-\pi, \pi]$

3) Physical Interpretation

When controlling drone attitude, we want the shortest rotation path:

- From yaw=170° to yaw=-170°: rotate 20° (NOT 340°)
- From pitch=3° to pitch=-3°: rotate 6° (straightforward)

Arguments:

angle_to: target angle (radians)
 angle_from: source angle (radians)

Returns:

difference: shortest angular distance (radians), wrapped to $[-\pi, \pi]$

Student code here

Calculating raw difference
`diff = angle_to - angle_from`

Standard wrapping to $[-\pi, \pi]$
`diff = (diff + np.pi) % (2.0 * np.pi) - np.pi`

Handling the special case near $\pm\pi$ boundary
 # If both angles are near the boundary but on opposite sides
`if abs(diff) < 0.5:`
 # From near $-\pi$ to near $+\pi$ (positive direction through wrap)
 if angle_to > 2.5 and angle_from < -2.5 and diff < 0:
 diff = -diff # Flip to positive
 # From near $+\pi$ to near $-\pi$ (negative direction through wrap)
 elif angle_from > 2.5 and angle_to < -2.5 and diff > 0:
 diff = -diff # Flip to negative

End student code

`return diff`

```
# Run tests for Part 6 helper functions
suite = unittest.TestSuite()
suite.addTest(TestTrajOptHelpers('test_angle_diff_small'))
suite.addTest(TestTrajOptHelpers('test_angle_diff_wrapping'))
suite.addTest(TestTrajOptHelpers('test_angle_diff_180_degrees'))

unittest.TextTestRunner().run(suite)
```

```
...
-----
Ran 3 tests in 0.004s
OK
<unittest.runner.TextTestResult run=3 errors=0 failures=0>
```

```
# TODO 18: Pack Decision Variables
def pack_decision_vars(states: np.ndarray, controls: np.ndarray, N: int) -> np.ndarray:
    """
    Pack states and controls into flat decision vector for optimization.

    1) Why This Matters:
    scipy.optimize.minimize requires a single flat vector of decision variables.
    We need to convert our structured trajectory (states and controls at each time step)
    into this flat format, and later unpack it back.

    2) Physical Interpretation
    - **States (first part):** All N+1 waypoints (start, intermediates, goal)
    - **Controls (second part):** All N control inputs (applied between waypoints)

    Arguments:
        states: (N+1) x 6 array [px, py, pz, yaw, pitch, roll]
        controls: N x 4 array [dyaw, dpitch, droll, thrust]
        N: number of time steps

    Returns:
        z: flat vector of length 10*N + 6
    """

##### Student code here #####
# Flattening all states the into a 1D array
# states is (N+1) x 6, so flatten gives us 6*(N+1) elements
states_flat = states.flatten()

# Flattening all controls into a 1D array
# controls is N x 4, so flatten gives us 4*N elements
controls_flat = controls.flatten()

# Concatenate: [all states, then all controls]
# Total length: 6*(N+1) + 4*N = 6N + 6 + 4N = 10N + 6
z = np.concatenate([states_flat, controls_flat])

##### End student code #####
return z
```

```
# TODO 19: Unpack Decision Variables
def unpack_decision_vars(z: np.ndarray, N: int) -> Tuple[np.ndarray, np.ndarray]:
    """
    Unpack flat decision vector into structured states and controls arrays.

    1) Why This Matters
    This is the inverse of `pack_decision_vars`. The optimizer works with flat vectors,
    but our cost/constraint functions need structured arrays. This function extracts
    the trajectory information from the flat optimization variable.

    2) Physical Interpretation
    We're converting from optimizer format (flat vector) back to trajectory format
    (time-indexed waypoints and control sequences).

    Arguments:
        z: flat vector of length 10*N + 6
        N: number of time steps

    Returns:
        states: (N+1) x 6 array [px, py, pz, yaw, pitch, roll]
        controls: N x 4 array [dyaw, dpitch, droll, thrust]
    """

##### Student code here #####
# Calculating how many elements belong to states
# We have (N+1) states, each with 6 values
```

```

num_state_elements = 6 * (N + 1)

# Extracting states from the beginning of z
states_flat = z[:num_state_elements]

# Reshaping into (N+1) x 6 matrix
states = states_flat.reshape(N + 1, 6)

# Extracting controls from the rest of z
controls_flat = z[num_state_elements:]

# Reshaping into N x 4 matrix
controls = controls_flat.reshape(N, 4)

##### End student code #####
return states, controls

```

```

# Run tests for Part 6 helper functions
suite = unittest.TestSuite()

suite.addTest(TestTrajOptHelpers('test_pack_decision_vars'))
suite.addTest(TestTrajOptHelpers('test_unpack_decision_vars'))
suite.addTest(TestTrajOptHelpers('test_pack_unpack_inverse'))

unittest.TextTestRunner().run(suite)

...
-----
Ran 3 tests in 0.007s

OK
<unittest.runner.TextTestResult run=3 errors=0 failures=0>

```

7.1 Defining the Cost Function

Now that we can pack and unpack our decision variables, we're ready to define the **cost function** that the optimizer will minimize. This is the heart of trajectory optimization!

What Are We Optimizing?

Remember, trajectory optimization is about finding the "best" trajectory through the hoops. But what does "best" mean? We need to mathematically define what makes one trajectory better than another.

Our total cost function has the form:

$$J(z) = J_{\text{thrust}}(z) + J_{\text{angular}}(z) + J_{\text{smooth}}(z) + J_{\text{gimbal}}(z)$$

Where:

- J_{thrust} : Penalizes deviation from hover thrust (energy efficiency)
- J_{angular} : Penalizes large angular velocities (aggressive maneuvers)
- J_{smooth} : Penalizes control changes between time steps (jerk)
- J_{gimbal} : Penalizes dangerous pitch angles near $\pm 90^\circ$ (gimbal lock)

Why Multiple Cost Terms?

Each term captures a different aspect of "good" flight:

1. Thrust Deviation (J_{thrust}):

- Encourages flight near hover condition ($T \approx 10$)
- Hovering is most energy-efficient
- Large thrust deviations means more battery consumption

2. Angular Velocity (J_{angular}):

- Penalizes rapid rotations (high $\Delta\psi, \Delta\theta, \Delta\phi$)
- Aggressive maneuvers stress motors and risk instability
- Encourages smooth, gentle turns

3. Control Smoothness (J_{smooth}):

- Penalizes sudden changes in control inputs
- Example: Going from T=5 to T=20 in one time step is jerky
- Smooth control changes → comfortable flight

4. Gimbal Lock (J_{gimbal}):

- Exponentially penalizes pitch near $\pm 90^\circ$
- At pitch = $\pm 90^\circ$, yaw and roll become indistinguishable (singularity!)
- Keeps drone away from dangerous attitude configurations

How the Optimizer Uses This

The `scipy.optimize.minimize` function will:

1. Start with an initial guess trajectory (start with the rrt traj)
2. Compute $J(z)$ for current trajectory
3. Try small variations to decrease J
4. Repeat until it finds a local minimum

The optimizer tries to make J as small as possible while satisfying all constraints (dynamics, boundary, collisions).

What You'll Implement

In the next series of TODOs, you will implement each cost term individually:

- **TODO 20:** `cost_function_thrust()` - Penalize thrust deviation from hover
- **TODO 21:** `cost_function_angular_velocity()` - Penalize large rotations
- **TODO 22:** `cost_function_smoothness()` - Penalize control jerk
- **TODO 23:** `cost_function_gimbal_lock()` - Penalize dangerous pitch
- **TODO 24:** `cost_function_integrated()` - Combine all terms with weights

After implementing each cost function, you'll run unit tests to verify correctness. The test cases check:

- Zero cost for ideal conditions (hover, no rotation, smooth controls)
- Positive cost for deviations from ideal
- Large penalties for dangerous configurations

Mathematical Notation Reminder

- N = number of time steps
- x_k = state at time k = $[p_x, p_y, p_z, \psi, \theta, \phi]_k$
- u_k = control at time k = $[\Delta\psi, \Delta\theta, \Delta\phi, T]_k$
- ψ (yaw), θ (pitch), ϕ (roll) in radians
- $T_{\text{hover}} = 10$ (hover thrust to counteract gravity)

Let's start implementing each cost component!

```
# Lets Define some Test Cases for the Cost Functions to check if they are working as expected!
class TestCostFunctions(unittest.TestCase):
    """Test all cost functions for trajectory optimization."""

    def test_cost_function_thrust_hover(self):
        """Test thrust cost at hover condition."""
        N = 2
        states = np.zeros((N+1, 6))
        # All thrust at hover (10 N)
        controls = np.array([[0, 0, 0, 10],
                            [0, 0, 0, 10]])

        cost = cost_function_thrust(states, controls, N, weight=0.1)
        # Thrust exactly at hover, so cost should be 0
        self.assertAlmostEqual(cost, 0.0, places=6)

    def test_cost_function_thrust_deviation(self):
        """Test thrust cost with deviation from hover."""
        N = 2
        states = np.zeros((N+1, 6))
        # Thrust deviates by +2 from hover (10 -> 12)
        controls = np.array([[0, 0, 0, 12],
                            [0, 0, 0, 12]])

        cost = cost_function_thrust(states, controls, N, weight=0.1)
```

```

# Cost = 0.1 * (2^2 + 2^2) = 0.1 * 8 = 0.8
self.assertAlmostEqual(cost, 0.8, places=6)

def test_cost_function_angular_zero(self):
    """Test angular cost with zero angular velocities."""
    N = 2
    states = np.zeros((N+1, 6))
    controls = np.array([[0, 0, 0, 10],
                        [0, 0, 0, 10]])

    cost = cost_function_angular(states, controls, N, weight=1.0)
    # All angular velocities zero, cost should be 0
    self.assertAlmostEqual(cost, 0.0, places=6)

def test_cost_function_angular_nonzero(self):
    """Test angular cost with non-zero angular velocities."""
    N = 2
    states = np.zeros((N+1, 6))
    # Angular velocities: [0.1, 0.2, 0.3]
    controls = np.array([[0.1, 0.2, 0.3, 10],
                        [0.1, 0.2, 0.3, 10]])

    cost = cost_function_angular(states, controls, N, weight=1.0)
    # Cost = 1.0 * ((0.1^2 + 0.2^2 + 0.3^2) + (0.1^2 + 0.2^2 + 0.3^2))
    #      = 1.0 * (0.14 + 0.14) = 0.28
    self.assertAlmostEqual(cost, 0.28, places=6)

def test_cost_function_smoothness_constant(self):
    """Test smoothness cost with constant controls."""
    N = 3
    states = np.zeros((N+1, 6))
    # Constant controls (no jerk)
    controls = np.array([[0.1, 0.1, 0.1, 10],
                        [0.1, 0.1, 0.1, 10],
                        [0.1, 0.1, 0.1, 10]])

    cost = cost_function_smoothness(states, controls, N, weight=5.0)
    # Controls don't change, so smoothness cost should be 0
    self.assertAlmostEqual(cost, 0.0, places=6)

def test_cost_function_smoothness_varying(self):
    """Test smoothness cost with varying controls."""
    N = 2
    states = np.zeros((N+1, 6))
    # Controls change from k=0 to k=1
    controls = np.array([[0, 0, 0, 10],
                        [0.1, 0.1, 0.1, 12]])

    cost = cost_function_smoothness(states, controls, N, weight=5.0)
    # Difference: [0.1, 0.1, 0.1, 2]
    # Cost = 5.0 * (0.1^2 + 0.1^2 + 0.1^2 + 2^2) = 5.0 * 4.03 = 20.15
    self.assertAlmostEqual(cost, 20.15, places=6)

def test_cost_function_gimbal_lock_safe(self):
    """Test gimbal lock penalty in safe range."""
    N = 2
    # Pitch within safe range (< 50°)
    states = np.array([[0, 0, 0, 0, 0.5, 0],      # pitch = 0.5 rad (~29°) - OK
                      [0, 0, 0, 0, 0.6, 0],      # pitch = 0.6 rad (~34°) - OK
                      [0, 0, 0, 0, 0.7, 0]])    # pitch = 0.7 rad (~40°) - OK
    controls = np.zeros((N, 4))

    cost = cost_function_gimbal_lock(states, controls, N)
    # All pitches within safe range, cost should be 0
    self.assertAlmostEqual(cost, 0.0, places=6)

def test_cost_function_gimbal_lock_danger(self):
    """Test gimbal lock penalty approaching singularity."""
    N = 2
    pitch_limit = 50 * np.pi / 180 # ~0.873 rad
    # Pitch exceeds safe range
    states = np.array([[0, 0, 0, 0, 1.0, 0],      # pitch = 1.0 rad (~57°) - DANGER
                      [0, 0, 0, 0, 0.5, 0],      # pitch = 0.5 rad - OK
                      [0, 0, 0, 0, -1.0, 0]])   # pitch = -1.0 rad - DANGER
    controls = np.zeros((N, 4))

    cost = cost_function_gimbal_lock(states, controls, N)

```

```

# Two states exceed limit
# Excess for pitch=1.0: 1.0 - 0.873 = 0.127
# Cost per violation: 1000 * 0.127^2 ≈ 16.13
# Total: 2 * 16.13 ≈ 32.26
self.assertTrue(cost > 30.0) # Should have significant penalty

def test_cost_function_integration(self):
    """Test integrated cost function combines all costs."""
    N = 2
    states = np.array([[0, 0, 0, 0, 0, 0],
                      [1, 1, 1, 0.1, 0.1, 0.1],
                      [2, 2, 2, 0.2, 0.2, 0.2]])
    controls = np.array([[0.1, 0.1, 0.1, 12],
                        [0.1, 0.1, 0.1, 12]])

    z = pack_decision_vars(states, controls, N)
    weights = {'thrust': 0.1, 'angular': 1.0, 'smoothness': 5.0}

    cost = cost_function_tuned(z, N, weights)

    # Should be sum of individual costs
    cost_thrust = cost_function_thrust(states, controls, N, weights['thrust'])
    cost_angular = cost_function_angular(states, controls, N, weights['angular'])
    cost_smooth = cost_function_smoothness(states, controls, N, weights['smoothness'])
    cost_gimbal = cost_function_gimbal_lock(states, controls, N)

    expected_cost = cost_thrust + cost_angular + cost_smooth + cost_gimbal
    self.assertAlmostEqual(cost, expected_cost, places=6)

```

```

# TODO 20: Cost Function - Thrust Deviation
def cost_function_thrust(states: np.ndarray, controls: np.ndarray, N: int,
                         weight: float = 0.1) -> float:
    """
    Compute thrust deviation cost (penalizes deviation from hover thrust).

```

Physical Interpretation

- Thrust = 10 N: hovering (zero cost contribution)
- Thrust = 15 N: climbing or accelerating (cost = w × 25)
- Thrust = 5 N: descending (cost = w × 25)

Higher thrust -> more power consumption -> higher cost.

Implementation Hints

Arguments:

```

states: (N+1) x 6 array (not used for this cost)
controls: N x 4 array [dyaw, dpitch, droll, thrust]
N: number of time steps
weight: cost weight (default 0.1)

```

Returns:

```

cost: scalar thrust deviation cost
"""

```

```
T_hover = 10.0
```

```
cost = 0.0
```

```
##### Student code here #####
```

```
# Looping through all time steps
```

```
for k in range(N):
```

```
    # Extracting thrust at time step k (4th element of control vector)
    thrust = controls[k, 3]
```

```
    # Computing deviation from hover thrust
    deviation = thrust - T_hover
```

```
    # Adding weighted squared deviation to total cost
    cost += weight * (deviation ** 2)
```

```
##### End student code #####
return cost
```

```

suite = unittest.TestSuite()
suite.addTest(TestCostFunctions('test_cost_function_thrust_hover'))

```

```

suite.addTest(TestCostFunctions('test_cost_function_thrust_deviations'))

unittest.TextTestRunner().run(suite)

..
-----
Ran 2 tests in 0.005s

OK
<unittest.runner.TextTestResult run=2 errors=0 failures=0>

```

```

# TODO 21: Cost Function - Angular Velocity
def cost_function_angular(states: np.ndarray, controls: np.ndarray, N: int,
                           weight: float = 1.0) -> float:
    """
    Compute angular velocity cost (penalizes large attitude changes).

    Physical Interpretation
    - Small angular changes (< 5°): minimal cost, smooth rotation
    - Large angular changes (> 20°): high cost, aggressive maneuver
    - Zero angular change: hovering in place (zero cost)
    Minimizing this encourages the drone to maintain stable orientation.

    Hint: Use numpy operations to compute squared norms efficiently.

    Arguments:
        states: (N+1) x 6 array (not used for this cost)
        controls: N x 4 array [dyaw, dpitch, droll, thrust]
        N: number of time steps
        weight: cost weight (default 1.0)

    Returns:
        cost: scalar angular velocity cost
    """
    cost = 0.0

##### Student code here #####
    # Loop through all time steps
    for k in range(N):
        # Extract angular velocities (first 3 elements of control vector)
        dyaw = controls[k, 0]
        dpitch = controls[k, 1]
        droll = controls[k, 2]

        # Add weighted sum of squared angular velocities
        cost += weight * (dyaw**2 + dpitch**2 + droll**2)

    ##### End student code #####
    return cost

```

```

suite = unittest.TestSuite()

suite.addTest(TestCostFunctions('test_cost_function_angular_zero'))
suite.addTest(TestCostFunctions('test_cost_function_angular_nonzero'))

unittest.TextTestRunner().run(suite)

..
-----
Ran 2 tests in 0.002s

OK
<unittest.runner.TextTestResult run=2 errors=0 failures=0>

```

```

# TODO 22: Cost Function - Control Smoothness
def cost_function_smoothness(states: np.ndarray, controls: np.ndarray, N: int,
                               weight: float = 5.0) -> float:
    """
    Compute control smoothness cost (jerk minimization).

    1) Why This Matters
    Jerk is the derivative of acceleration (third derivative of position).
    High jerk causes:

```

- Mechanical wear on actuators
 - Uncomfortable flight dynamics - especially important if we had human passengers
 - Tracking errors in control systems
- Smooth control sequences -> smooth trajectories -> efficient flight.

2) Physical Interpretation

- Constant controls over time: zero smoothness cost (ideal)
- Gradually changing controls: low cost (acceptable)
- Sudden control jumps: high cost (penalized)

This encourages smooth transitions, not abrupt changes.

Hint: Use numpy operations to compute differences and squared norms efficiently.

Arguments:

```
states: (N+1) x 6 array (not used for this cost)
controls: N x 4 array [dyaw, dpitch, droll, thrust]
N: number of time steps
weight: cost weight (default 5.0)
```

Returns:

```
cost: scalar smoothness cost
```

```
"""
```

```
if N <= 1:
    return 0.0
```

```
cost = 0.0
```

```
##### Student code here #####
```

```
# Loop through consecutive time steps
```

```
for k in range(N - 1):
    # Computing the difference between consecutive controls
    control_diff = controls[k + 1] - controls[k]
```

```
# Adding weighted sum of squared differences for all 4 control elements
cost += weight * np.sum(control_diff ** 2)
```

```
##### End student code #####
return cost
```

```
suite = unittest.TestSuite()

suite.addTest(TestCostFunctions('test_cost_function_smoothness_constant'))
suite.addTest(TestCostFunctions('test_cost_function_smoothness_varying'))

unittest.TextTestRunner().run(suite)
```

```
..
-----
Ran 2 tests in 0.005s
```

```
OK
<unittest.runner.TextTestResult run=2 errors=0 failures=0>
```

```
# TODO 23: Cost Function - Gimbal Lock Penalty
def cost_function_gimbal_lock(states: np.ndarray, controls: np.ndarray, N: int) -> float:
    """
    Soft penalty to avoid gimbal lock singularity.
    
```

1) Why This Matters

Gimbal lock occurs when pitch = $\pm 90^\circ$ in Euler angles (ZYX convention).

At this singularity:

- Yaw and roll become indistinguishable
- Rotation matrix loses a degree of freedom
- Numerical instability in euler_from_rotation_matrix

We add a soft penalty when $|pitch| > 50^\circ$ to keep the optimizer away from this region.

2) Physical Interpretation

- Pitch in $[-50^\circ, 50^\circ]$: normal operation, zero penalty
- Pitch approaching $\pm 90^\circ$: rapidly increasing penalty
- Prevents optimizer from exploring gimbal lock region

This is a **soft constraint** (high cost) rather than a hard constraint.

Arguments:

```
states: (N+1) x 6 array [px, py, pz, yaw, pitch, roll]
```

```

controls: N x 4 array (not used for this cost)
N: number of time steps

Returns:
    cost: scalar gimbal lock penalty
"""

cost = 0.0

##### Student code here #####
# Define pitch limit in radians (50 degrees)
pitch_limit = 50 * np.pi / 180 # ~0.873 rad
penalty_weight = 1000.0

# Looping through all states (N+1 states for N time steps)
for k in range(N + 1):
    # Extracting pitch angle (index 4 in state vector)
    pitch = states[k, 4]

    # Checking if absolute pitch exceeds the safe limit
    if np.abs(pitch) > pitch_limit:
        # Computing how much pitch exceeds the limit
        excess = np.abs(pitch) - pitch_limit

        # Adding quadratic penalty for the excess
        cost += penalty_weight * (excess ** 2)

#####
# End student code #####
return cost

```

```

suite = unittest.TestSuite()

suite.addTest(TestCostFunctions('test_cost_function_gimbal_lock_safe'))
suite.addTest(TestCostFunctions('test_cost_function_gimbal_lock_danger'))

unittest.TextTestRunner().run(suite)

..
-----
Ran 2 tests in 0.005s

OK
<unittest.runner.TextTestResult run=2 errors=0 failures=0>

```

```

# Now we can define the integrated cost function that combines all the individual cost components with tunable weights!
def cost_function_tuned(z: np.ndarray, N: int, weights: Dict[str, float]) -> float:
"""
    Integrated cost function combining all cost components.

```

1) Why This Matters

This is the **objective function** that the optimizer minimizes. It combines all the cost components (thrust, angular, smoothness, gimbal lock) with tunable weights to achieve desired flight characteristics.

2) Physical Interpretation

Weight Tuning Strategy:

- **Racing mode** (fast, aggressive):
 - thrust: 0.05 (allow large thrust changes)
 - angular: 0.5 (allow aggressive turns)
 - smoothness: 5.0 (moderate smoothing)
- **Normal mode** (smooth, efficient):
 - thrust: 0.1 (prefer hover)
 - angular: 1.0 (smooth turns)
 - smoothness: 10.0 (high smoothing)

Arguments:

```

z: decision variables (flat vector of length 10*N + 6)
N: number of time steps
weights: dict with keys 'thrust', 'angular', 'smoothness'

```

Returns:

```

cost: scalar total cost
"""

```

```

states, controls = unpack_decision_vars(z, N)

cost = 0.0

# Add thrust cost
cost += cost_function_thrust(states, controls, N, weights['thrust'])

# Add angular velocity cost
cost += cost_function_angular(states, controls, N, weights['angular'])

# Add smoothness cost
cost += cost_function_smoothness(states, controls, N, weights['smoothness'])

# Add gimbal lock penalty
cost += cost_function_gimbal_lock(states, controls, N)

return cost

```

```

suite = unittest.TestSuite()

suite.addTest(TestCostFunctions('test_cost_function_integration'))

unittest.TextTestRunner().run(suite)

.

-----
Ran 1 test in 0.003s

OK
<unittest.runner.TextTestResult run=1 errors=0 failures=0>

```

✓ 7.2 Defining Constraints: Making Physics Mandatory

We've defined our cost function $J(z)$ - but cost alone is not enough! Without constraints, the optimizer might find "solutions" that:

- Teleport between waypoints (violate physics)
- Start at the wrong location
- Miss the goal position
- Skip hoops entirely

Constraints make physically impossible trajectories mathematically impossible.

What Are Constraints?

Constraints are **equations or inequalities** that the optimizer MUST satisfy. There are two types:

1. **Equality Constraints:** $c(z) = 0$ (must be exactly satisfied)
 - Examples: "Start at position (1, 3, 8)", "Obey dynamics", "Pass through hoop center"
2. **Inequality Constraints:** $c(z) \leq 0$ (must be non-positive)
 - Examples: "Stay at least 0.5m from obstacles", "Thrust ≤ 20 N"

The optimizer searches for trajectories z that:

- **Minimize** $J(z)$ (cost)
- **While satisfying** all constraints

Our Constraint Types

We implement three types of equality constraints:

1. Dynamics Constraints (TODO 24) - THE MOST IMPORTANT

These ensure the trajectory follows the **laws of physics**. For each time step $k = 0, \dots, N - 1$:

Position dynamics (3 equations):

$$\mathbf{p}_{k+1} = \mathbf{p}_k + \mathbf{v}_{\text{terminal}}(\psi_k, \theta_k, \phi_k, T_k) \cdot \Delta t$$

Attitude dynamics (3 equations):

$$\begin{aligned}\psi_{k+1} &= \psi_k + \Delta\psi_k \\ \theta_{k+1} &= \theta_k + \Delta\theta_k \\ \phi_{k+1} &= \phi_k + \Delta\phi_k\end{aligned}$$

Total: $6N$ constraints (6 per time step)

2. Boundary Constraints (TODO 25)

These fix the **start and goal** of the trajectory:

Start constraints (6 equations):

$$\mathbf{x}_0 = \mathbf{x}_{\text{start}} \Rightarrow \begin{bmatrix} p_x \\ p_y \\ p_z \\ \psi \\ \theta \\ \phi \end{bmatrix}_0 = \begin{bmatrix} 1 \\ 3 \\ 8 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Goal position constraints (3 equations):

$$\mathbf{p}_N = \mathbf{p}_{\text{goal}} \Rightarrow \begin{bmatrix} p_x \\ p_y \\ p_z \end{bmatrix}_N = \begin{bmatrix} x_{\text{final}} \\ y_{\text{final}} \\ z_{\text{final}} \end{bmatrix}$$

Hoop waypoint constraints (3H equations, where H = number of hoops): For each hoop h , at designated knot index k_h :

$$\mathbf{p}_{k_h} = \mathbf{p}_{\text{hoop}_h}$$

Total: $9 + 3H$ constraints (9 boundary + 3 per hoop)

Why this matters: These ensure we start at the right place, end at the right place, and pass through each hoop center.

3. Collision Constraints (We define it for you)

These ensure the drone maintains a **safety margin** from obstacles:

For sphere obstacles:

$$c_{\text{sphere}}(z) = (r_{\text{obs}} + r_{\text{safety}}) - \|\mathbf{p}_k - \mathbf{c}_{\text{obs}}\|_2 \leq 0$$

Equivalently: $\|\mathbf{p}_k - \mathbf{c}_{\text{obs}}\|_2 \geq r_{\text{obs}} + r_{\text{safety}}$

Total: Depends on number of obstacles (this is an inequality constraint)

How Constraints Work Mathematically

For equality constraints $c(z) = 0$, we return the **violation**:

$$\text{violation} = \text{actual_value} - \text{expected_value}$$

The optimizer adjusts z until all violations are (approximately) zero.

Constraint Dimensions Summary

For a trajectory with N time steps and H hoops:

Constraint Type	Count	Dimension	Type
Dynamics	$6N$	\mathbb{R}^{6N}	Equality
Boundary (start)	6	\mathbb{R}^6	Equality
Boundary (goal)	3	\mathbb{R}^3	Equality
Hoops	$3H$	\mathbb{R}^{3H}	Equality
Total Equality	$6N + 9 + 3H$	-	-

Example: For $N = 15$ time steps and $H = 4$ hoops:

- Decision variables: $10(15) + 6 = 156$
- Dynamics constraints: $6(15) = 90$
- Boundary constraints: 9
- Hoop constraints: $3(4) = 12$
- **Total constraints:** $90 + 9 + 12 = 111$ **equations**

The optimizer must find 156 variables satisfying 111 equations while minimizing cost!

What You'll Implement

- **TODO 24:** `dynamics_constraints_robust()` - Enforce physics (terminal velocity model)
- **TODO 25:** `boundary_constraints_with_hoops()` - Fix start, goal, and hoop passages

After implementing, you'll test each constraint function to verify:

- Zero violations for valid trajectories
- Non-zero violations for invalid trajectories
- Correct dimensions (number of constraint equations)

The Big Picture: Cost + Constraints

Our complete optimization problem is:

$\min_{z \in \mathbb{R}^{10N+6}} J(z)$	(minimize cost)
subject to:	
$\mathbf{c}_{\text{dyn}}(z) = \mathbf{0}$	(obey physics)
$\mathbf{c}_{\text{boundary}}(z) = \mathbf{0}$	(correct start/goal)
$\mathbf{c}_{\text{hoop}}(z) = \mathbf{0}$	(pass through hoops)
$\mathbf{c}_{\text{collision}}(z) \leq \mathbf{0}$	(avoid obstacles)

The optimizer (`scipy.optimize.minimize` with SLSQP method) uses gradient information to efficiently search the space of feasible trajectories for the one with lowest cost.

Let's implement the constraints!

```
# Lets define some Test Cases for the Constraint Functions to check if they are working as expected!
class TestConstraints(unittest.TestCase):
    """Test dynamics and boundary constraint functions."""

    def test_dynamics_constraints_hover(self):
        """Test dynamics constraints for hovering (stationary) trajectory."""
        N = 2
        dt = 0.1

        # Hovering: same position, zero attitude, hover thrust
        states = np.array([[5, 5, 5, 0, 0, 0],
                           [5, 5, 5, 0, 0, 0],
                           [5, 5, 5, 0, 0, 0]])
        controls = np.array([[0, 0, 0, 10],      # Zero angular changes, hover thrust
                            [0, 0, 0, 10]])

        z = pack_decision_vars(states, controls, N)
        violations = dynamics_constraints_robust(z, N, dt)

        # Check shape: should be 6*N = 12 constraints
        self.assertEqual(violations.shape[0], 6 * N)

        # Hovering should nearly satisfy dynamics (small violations due to drag)
        # Position should change slightly due to terminal velocity
        # But violations should be small
        self.assertTrue(np.max(np.abs(violations)) < 1.0)

    def test_dynamics_constraints_forward_flight(self):
        """Test dynamics constraints for forward flight."""
        N = 2
        dt = 0.1

        # Forward flight: moving in +x direction with pitch
        states = np.array([[0, 0, 5, 0, 0.2, 0],  # pitch forward 0.2 rad
                           [1, 0, 5, 0, 0.2, 0],  # moved 1m forward
                           [2, 0, 5, 0, 0.2, 0]]) # moved 2m total
        controls = np.array([[0, 0, 0, 15],          # More thrust for forward flight
                            [0, 0, 0, 15]])

        z = pack_decision_vars(states, controls, N)
        violations = dynamics_constraints_robust(z, N, dt)

        # Check shape
        self.assertEqual(violations.shape[0], 6 * N)

        # Violations won't be zero (we didn't compute exact dynamics)
        # But they should exist (we're testing the function runs)
        self.assertTrue(isinstance(violations, np.ndarray))

    def test_boundary_constraints_start(self):
        """Test boundary constraints enforce start pose."""
        N = 2
        start_pose = gtsam.Pose3(gtsam.Rot3.Ypr(0.1, 0.2, 0.3),
                                  gtsam.Point3(1, 2, 3))
```

```

goal_position = np.array([8, 9, 10])

# States that match start pose
states = np.array([[1, 2, 3, 0.1, 0.2, 0.3], # Matches start
                  [4, 5, 6, 0.1, 0.2, 0.3],
                  [8, 9, 10, 0.1, 0.2, 0.3]]) # Matches goal position
controls = np.zeros((N, 4))

z = pack_decision_vars(states, controls, N)
violations = boundary_constraints_robust(z, N, start_pose, goal_position, [])

# Check shape: 6 (start) + 3 (goal) + 0 (hoops) = 9
self.assertEqual(violations.shape[0], 9)

# First 6 violations (start pose) should be near zero
self.assertTrue(np.max(np.abs(violations[:6])) < 0.1)

# Last 3 violations (goal position) should be near zero
self.assertTrue(np.max(np.abs(violations[6:9])) < 0.1)

def test_boundary_constraints_with_hoops(self):
    """Test boundary constraints with hoop waypoints."""
    N = 10
    start_pose = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(0, 0, 0))
    goal_position = np.array([10, 0, 0])

    # Create states with 2 hoops at specific locations
    states = np.zeros((N+1, 6))
    states[:, 0] = np.linspace(0, 10, N+1) # x from 0 to 10
    # Make sure some states pass near hoop positions
    states[3, :] = [3, 5, 5, 0, 0, 0] # Near hoop 1
    states[7, :] = [7, 8, 8, 0, 0, 0] # Near hoop 2
    states[N, :3] = goal_position # Goal position

    controls = np.zeros((N, 4))

    # Define 2 hoops
    hoop1 = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(3, 5, 5))
    hoop2 = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(7, 8, 8))
    hoops = [hoop1, hoop2]

    z = pack_decision_vars(states, controls, N)
    violations = boundary_constraints_robust(z, N, start_pose, goal_position, hoops)

    # Check shape: 6 (start) + 3 (goal) + 3*2 (hoops) = 15
    self.assertEqual(violations.shape[0], 15)

    # All violations should exist (function runs correctly)
    self.assertTrue(isinstance(violations, np.ndarray))

def test_boundary_constraints_dimensions(self):
    """Test boundary constraints have correct dimensions for various hoop counts."""
    N = 5
    start_pose = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(0, 0, 0))
    goal_position = np.array([5, 5, 5])

    states = np.random.randn(N+1, 6)
    states[0, :] = [0, 0, 0, 0, 0, 0]
    states[N, :3] = goal_position
    controls = np.zeros((N, 4))
    z = pack_decision_vars(states, controls, N)

    # Test with 0 hoops
    violations = boundary_constraints_robust(z, N, start_pose, goal_position, [])
    self.assertEqual(violations.shape[0], 9) # 6 + 3 + 0

    # Test with 1 hoop
    hoop1 = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(2, 2, 2))
    violations = boundary_constraints_robust(z, N, start_pose, goal_position, [hoop1])
    self.assertEqual(violations.shape[0], 12) # 6 + 3 + 3

    # Test with 3 hoops
    hoop2 = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(3, 3, 3))
    hoop3 = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(4, 4, 4))
    violations = boundary_constraints_robust(z, N, start_pose, goal_position, [hoop1, hoop2, hoop3])
    self.assertEqual(violations.shape[0], 18) # 6 + 3 + 9

```

```

# TODO 24: Dynamics Constraints - Terminal Velocity Model
def dynamics_constraints_robust(z: np.ndarray, N: int, dt: float) -> np.ndarray:
    """
    Enforce physics-based dynamics constraints (terminal velocity model).
    """

    Physical Interpretation
    Think of this as "consistency checking":
    - **Position**: "If I apply thrust T with attitude ( $\psi, \theta, \phi$ ), do I end up at  $p_{k+1}$ ?"
    - **Attitude**: "If I apply angular control  $\Delta\eta$ , do I get attitude  $\eta_{k+1}$ ?"

    The optimizer searches for (states, controls) where physics is respected.

    Hint: Use the 'compute_terminal_velocity' function from TODO 8 to get terminal velocity.

    Arguments:
        z: decision variables (flat vector of length 10*N + 6)
        N: number of time steps
        dt: time step duration (seconds, typically 0.1)

    Returns:
        violations: array of length 6N (3 position + 3 attitude per timestep)
    """
    violations = []
    mass = 1.0 # kg
    g = 10.0 # m/s^2 (gravity)

    ##### Student code here #####
    # Unpack decision variables
    states, controls = unpack_decision_vars(z, N)

    # Looping through each time step
    for k in range(N):
        # Extracting the current state and control
        px, py, pz = states[k, 0], states[k, 1], states[k, 2]
        yaw, pitch, roll = states[k, 3], states[k, 4], states[k, 5]

        dyaw, dpitch, droll, thrust = controls[k, 0], controls[k, 1], controls[k, 2], controls[k, 3]

        # Extracting the next state
        px_next, py_next, pz_next = states[k+1, 0], states[k+1, 1], states[k+1, 2]
        yaw_next, pitch_next, roll_next = states[k+1, 3], states[k+1, 4], states[k+1, 5]

        # Computing the force in navigation frame from thrust and attitude
        R = gtsam.Rot3.Ypr(yaw, pitch, roll)
        thrust_body = gtsam.Point3(0, 0, thrust) # Thrust in body frame (z-axis)
        thrust_nav = R.rotate(thrust_body) # Rotate to navigation frame

        # Adding the gravity (pointing down in nav frame)
        gravity = gtsam.Point3(0, 0, -mass * g)
        total_force = thrust_nav + gravity

        # Computing terminal velocity
        v_terminal = compute_terminal_velocity(total_force)

        # Position dynamics:  $p_{k+1} = p_k + v_{terminal} * dt$ 
        px_expected = px + v_terminal[0] * dt
        py_expected = py + v_terminal[1] * dt
        pz_expected = pz + v_terminal[2] * dt

        # Attitude dynamics:  $\eta_{k+1} = \eta_k + \Delta\eta$ 
        yaw_expected = yaw + dyaw
        pitch_expected = pitch + dpitch
        roll_expected = roll + droll

        # Computing violations (actual - expected)
        violations.append(px_next - px_expected)
        violations.append(py_next - py_expected)
        violations.append(pz_next - pz_expected)
        violations.append(yaw_next - yaw_expected)
        violations.append(pitch_next - pitch_expected)
        violations.append(roll_next - roll_expected)

    #### End student code ####

```

```

return np.array(violations)

suite = unittest.TestSuite()
suite.addTest(TestConstraints('test_dynamics_constraints_hover'))
suite.addTest(TestConstraints('test_dynamics_constraints_forward_flight'))

unittest.TextTestRunner().run(suite)

..
-----
Ran 2 tests in 0.005s

OK
<unittest.runner.TextTestResult run=2 errors=0 failures=0>

# TODO 25: Boundary Constraints - Start, Goal, and Hoops
def boundary_constraints_robust(z: np.ndarray, N: int, start_pose: gtsam.Pose3,
                                  goal_position: np.ndarray, hoops: List[gtsam.Pose3]) -> np.ndarray:
    """
    Enforce start, goal, and hoop waypoint constraints.

    Physical Interpretation
    - Start constraint: "The trajectory must begin exactly where the RRT path starts"
    - Goal constraint: "The trajectory must reach the target position"
    - Hoop constraints: "The trajectory must pass through the center of each hoop"
    The optimizer can adjust intermediate waypoints, but these boundary points are fixed.

    Hints:
    - Use 'helpers.euler_from_rotation_matrix_safe' to extract yaw, pitch, roll from rotation matrix.
    - Use 'helpers.find_hoop_indices_robust' to find which trajectory indices correspond to each hoop.
    - Use 'angle_diff' to compute angular differences (e.g., for yaw, pitch, roll).
    - Remember that start_pose is a full Pose3 (position + orientation), while goal_position is just a 3D point.

    Arguments:
        z: decision variables (flat vector)
        N: number of time steps
        start_pose: initial gtsam.Pose3 (position + orientation)
        goal_position: final position as numpy array [x, y, z]
        hoops: list of gtsam.Pose3 representing hoop positions

    Returns:
        violations: array of length (9 + 3*len(hoops))
    """
    violations = []

    ##### Student code here #####
    # Unpack decision variables
    states, controls = unpack_decision_vars(z, N)

    # 1. Starting constraints (6 constraints: 3 position + 3 attitude)
    start_position = start_pose.translation()
    start_rotation = start_pose.rotation()
    start_yaw, start_pitch, start_roll = helpers.euler_from_rotation_matrix_safe(start_rotation.matrix())

    # Extracting the actual initial state
    actual_start_pos = states[0, :3]
    actual_start_yaw = states[0, 3]
    actual_start_pitch = states[0, 4]
    actual_start_roll = states[0, 5]

    # Position violations (3)
    violations.append(actual_start_pos[0] - start_position[0])
    violations.append(actual_start_pos[1] - start_position[1])
    violations.append(actual_start_pos[2] - start_position[2])

    # Attitude violations (3) - use angle_diff for angular differences
    violations.append(angle_diff(actual_start_yaw, start_yaw))
    violations.append(angle_diff(actual_start_pitch, start_pitch))
    violations.append(angle_diff(actual_start_roll, start_roll))

    # 2. Goal constraints (3 constraints: position only)
    actual_goal_pos = states[N, :3]

```

```

violations.append(actual_goal_pos[0] - goal_position[0])
violations.append(actual_goal_pos[1] - goal_position[1])
violations.append(actual_goal_pos[2] - goal_position[2])

# 3. Hoop constraints (3 constraints per hoop)
if len(hoops) > 0:
    # Extracting hoop positions from Pose3 objects
    hoop_positions = [hoop.translation() for hoop in hoops]

    # Finding which trajectory indices correspond to each hoop
    hoop_indices = helpers.find_hoop_indices_robust(states[:, :3], hoop_positions, N)

    # For each hoop, add position constraints
    for hoop_idx, hoop in enumerate(hoops):
        # Get the trajectory index for this hoop
        traj_idx = hoop_indices[hoop_idx]

        # Getting hoop center position
        hoop_position = hoop.translation()

        # Getting actual position at this trajectory index
        actual_pos = states[traj_idx, :3]

        # Adding position violations for this hoop
        violations.append(actual_pos[0] - hoop_position[0])
        violations.append(actual_pos[1] - hoop_position[1])
        violations.append(actual_pos[2] - hoop_position[2])

##### End student code #####
return np.array(violations)

```

```

suite = unittest.TestSuite()
suite.addTest(TestConstraints('test_boundary_constraints_start'))
suite.addTest(TestConstraints('test_boundary_constraints_with_hoops'))
suite.addTest(TestConstraints('test_boundary_constraints_dimensions'))

unittest.TextTestRunner().run(suite)

```

```

...
-----
Ran 3 tests in 0.006s

OK
<unittest.runner.TextTestResult run=3 errors=0 failures=0>

```

```

import numpy as np
# TODO 26 (Done by TAs)
# Now we will define the collision constraints function!
# Look into it to learn how collision avoidance can be implemented for spherical objects in trajectory optimization.

```

```

def collision_constraints_optimized(z: np.ndarray, N: int, obstacles: List,
                                      subsample: int = 2) -> np.ndarray:
    """
    Enforce collision avoidance with obstacles (inequality constraints).
    This is challenging because:
    - Inequality constraints are harder to satisfy than equality constraints
    - May make optimization slower or fail to converge
    - Requires careful tuning of safety margins
    """

```

Physical Interpretation

- **Safety margin (0.5m)**: Additional clearance beyond obstacle radius
- **Subsampling**: Check every Nth knot point (e.g., every 2nd) for efficiency
- **Sphere obstacles**: Simple distance check
- **Box obstacles**: More complex (check if inside, compute distance to faces)

Arguments:

`z`: decision variables
`N`: number of time steps
`obstacles`: list of obstacle objects (`SphereObstacle` or `BoxObstacle`)
`subsample`: check every 'subsample' knots (e.g., 2 = every other)

Returns:

`violations`: array of inequality constraint violations
Note: This returns violations where:
- Positive = collision detected

```

- Negative = safe
- The optimizer will try to make all violations ≤ 0
"""

states, controls = unpack_decision_vars(z, N)
violations = []
safety_margin = 0.5

for k in range(0, N + 1, subsample):
    px, py, pz = states[k, 0:3]
    # point = gtsam.Point3(px, py, pz)

    for obs in obstacles:
        if isinstance(obs, helpers.SphereObstacle):
            # Vectorized distance computation
            dist = np.linalg.norm(np.array([px, py, pz]) - np.array(obs.center))
            violations.append(obs.radius + safety_margin - dist)

return np.array(violations)

```

Now Lets initialize from RRT Path to intial guess for Optimization. We are writing this for you!!

```

# Initialization from RRT Path
def initialize_from_rrt_robust(rrt_path: List[gtsam.Pose3], N: int, dt: float,
                                start_pose: gtsam.Pose3) -> np.ndarray:
    """
    Convert RRT path to initial guess for optimization.

    1) Why This Matters
    A good initial guess is **critical** for nonlinear optimization convergence:
    - Bad initialization → optimizer gets stuck in local minimum
    - Good initialization → fast convergence to global optimum
    - RRT provides feasible (collision-free) initialization

    2) Physical Interpretation
    We're resampling the RRT path (which may have 50-200 waypoints) down to N+1
    knot points (typically 15-30) for optimization. Linear interpolation provides
    smooth transitions.

    Arguments:
        rrt_path: list of gtsam.Pose3 from RRT
        N: number of time steps for optimization
        dt: time step duration (not used but included for future extensions)
        start_pose: initial pose (for validation)

    Returns:
        z_init: initial decision vector of length 10*N + 6
    """

    path_length = len(rrt_path)
    states_init = np.zeros((N + 1, 6))

    # Resample RRT path to N+1 knot points
    for i in range(N + 1):
        # Linear interpolation index
        idx_float = i * (path_length - 1) / N
        idx_low = int(np.floor(idx_float))
        idx_high = min(int(np.ceil(idx_float)), path_length - 1)
        alpha = idx_float - idx_low

        # Interpolate position
        pos_low = rrt_path[idx_low].translation()
        if idx_high > idx_low:
            pos_high = rrt_path[idx_high].translation()
            pos = pos_low + alpha * (pos_high - pos_low)
        else:
            pos = pos_low

        # Interpolate attitude (linear on Euler angles with wrapping)
        R_low = rrt_path[idx_low].rotation().matrix()
        yaw_low, pitch_low, roll_low = helpers.euler_from_rotation_matrix_safe(R_low)

        if idx_high > idx_low:
            R_high = rrt_path[idx_high].rotation().matrix()
            yaw_high, pitch_high, roll_high = helpers.euler_from_rotation_matrix_safe(R_high)

```

```

        # Handle angle wrapping
        yaw = yaw_low + alpha * angle_diff(yaw_high, yaw_low)
        pitch = pitch_low + alpha * angle_diff(pitch_high, pitch_low)
        roll = roll_low + alpha * angle_diff(roll_high, roll_low)
    else:
        yaw, pitch, roll = yaw_low, pitch_low, roll_low

    states_init[i, :] = [pos[0], pos[1], pos[2], yaw, pitch, roll]

    # Initialize controls: compute from state differences
    controls_init = np.zeros((N, 4))
    for k in range(N):
        # Angle changes
        controls_init[k, 0] = angle_diff(states_init[k + 1, 3], states_init[k, 3]) # dyaw
        controls_init[k, 1] = angle_diff(states_init[k + 1, 4], states_init[k, 4]) # dpitch
        controls_init[k, 2] = angle_diff(states_init[k + 1, 5], states_init[k, 5]) # droll

        # Thrust: start with hover thrust
        controls_init[k, 3] = 10.0 # Newtons

    # Clamp controls to bounds (vectorized)
    deg_to_rad = np.pi / 180
    controls_init[:, 0:3] = np.clip(controls_init[:, 0:3], -10 * deg_to_rad, 10 * deg_to_rad)
    controls_init[:, 3] = np.clip(controls_init[:, 3], 5, 20)

    return pack_decision_vars(states_init, controls_init, N)

```

▼ 7.3 Putting It All Together: The Optimizer!

Okay, take a deep breath. You've just implemented a LOT of math:

- Cost functions (thrust, angular, smoothness, gimbal lock)
- Dynamics constraints (physics!)
- Boundary constraints (start, goal, hoops)
- Collision constraints (obstacles, optional)
- Helper functions (pack, unpack, angle wrapping)

Now comes the **REALLY cool part** - we're going to unleash `scipy.optimize.minimize` to find the optimal trajectory!

What Does the Optimizer Actually Do?

Remember this beast of an optimization problem?

$\min_{z \in \mathbb{R}^{10N+6}}$	$J(z)$	(your cost functions)
subject to:	$\mathbf{c}_{\text{dyn}}(z) = \mathbf{0}$	(your dynamics constraints)
	$\mathbf{c}_{\text{boundary}}(z) = \mathbf{0}$	(your boundary constraints)
	$\mathbf{c}_{\text{collision}}(z) \leq \mathbf{0}$	(your collision constraints)

The optimizer is going to:

1. **Start** with your RRT path as an initial guess (156 variables for N=15!)
2. **Compute** the cost $J(z)$ and all constraint violations
3. **Use gradients** (calculus!) to figure out which direction to move z
4. **Take a step** that decreases cost while respecting constraints
5. **Repeat** steps 2-4 until convergence (typically 20-50 iterations for a free environment but can get complex for a obstacle course)

This is **nonlinear constrained optimization** - one of the most powerful tools in robotics!

The Magic of SLSQP (Sequential Least Squares Programming)

We use `scipy.optimize.minimize` with the `SLSQP` method. Why SLSQP?

- **Handles equality AND inequality constraints** (most optimizers can't do both!)
- **Gradient-based** = fast convergence (beats genetic algorithms by 100x)
- **Battle-tested** = used in aerospace, robotics

How it works (simplified):

1. Linearize cost and constraints around current point
2. Solve a quadratic program (QP) to get search direction

3. Line search to find best step size
4. Update decision variables
5. Check convergence (gradient norm, constraint violations)

What Makes Optimization Succeed or Fail?

Success factors:

- Good RRT initialization (feasible path)
- Correct constraint dimensions (you tested these!)
- Balanced cost weights (not too large, not too small)
- Reasonable N value (15-30 knots works well)

Failure modes:

- Bad initialization (path goes through obstacle)
- Constraint dimension mismatch (crashes immediately)
- Poorly scaled costs (one term dominates)
- Too many knots (N > 50 becomes slow)
- Infeasible problem (impossible to reach goal while avoiding obstacles)

When optimization fails, we have a fallback strategy:

1. Try SLSQP first (fast, but sensitive)
2. If SLSQP fails, try `trust-constr` (slower, more robust)
3. If both fail, return RRT path (still collision-free!)

The Complete Optimization Workflow

Here's what happens when you call `optimize_trajectory()`:

```
optimized_path, success, info = optimize_trajectory(
    rrt_path=rrt_path,           # Your RRT path from Part 5
    start_pose=start_pose,       # Starting pose
    goal_position=goal_pos,     # Goal position
    hoops=hoops,                # List of hoop poses
    obstacles=obstacles,         # List of obstacles
    N=20,                      # Number of knot points
    dt=0.1,                     # Time step
    weights={'thrust': 0.1, 'angular': 1.0, 'smoothness': 5.0}
)
```

Step-by-step what happens inside:

1. Initialize
 - Resample RRT path to N+1 knot points
 - Compute initial controls
 - Pack into decision vector z_0

2. Define objective function for scipy

```
def objective(z):
    states, controls = unpack_decision_vars(z, N)
    cost = cost_function_integrated(states, controls, N, weights)
    return cost
```

3. Define equality constraints for scipy

```
eq_constraints = [
    {'type': 'eq', 'fun': lambda z: dynamics_constraints_robust(z, N, dt)},
    {'type': 'eq', 'fun': lambda z: boundary_constraints_robust(z, N, start, goal, hoops)}
]
```

4. Define inequality constraints (if obstacles present)

```
ineq_constraints = [
    {'type': 'ineq', 'fun': lambda z: -collision_constraints_optimized(z, N, obstacles)}
```

```
[]
# Note: scipy wants g(z) >= 0, but we return violations where positive = bad
# So we negate: -collision_constraints makes negative violations become positive (good)
```

5. Call `scipy.optimize.minimize`

```
result = scipy.optimize.minimize(
    objective,
    z_init,
    method='SLSQP',
    constraints=eq_constraints + ineq_constraints,
    options={'maxiter': 400, 'ftol': 1e-5}
)
```

6. Extract optimized trajectory

```
z_opt = result.x
states_opt, controls_opt = unpack_decision_vars(z_opt, N)
# Convert states to Pose3 list for visualization
```

7. Validate solution

- Check constraint violations < 0.01
- Check cost is reasonable
- Check path is collision-free

What You've Built: A Production-Grade Trajectory Optimizer!

Let's put this in perspective. You've implemented the **same core algorithm** used by:

- **Drone racing companies** (Skydio, DJI)
- **SpaceX** (Falcon 9 landing trajectories)
- **Self-driving cars** (motion planning)
- **Humanoid robots** (Boston Dynamics Atlas)
- **Aircraft autopilots** (Boeing, Airbus)

The only differences are:

- More complex dynamics models (yours is 6-DOF, theirs might be 12-DOF)
- More sophisticated cost functions (fuel optimization, passenger comfort)
- Real-time implementation (MPC: Model Predictive Control)

But the core math? Identical. Direct transcription + constrained optimization = industry standard.

Your TODOs Are Done - Now Let's FLY!

You've completed the TODOs. All the hard math is implemented. Now we get to the fun part:

In the next cells, you'll:

- Run optimization on simple paths (no hoops)
- Run optimization on racing paths (with hoops!)
- Compare RRT vs optimized paths side-by-side
- Visualize velocity and acceleration profiles
- See your drone smoothly flying through hoops
- Challenge yourself with obstacles

The optimizer functions (`optimize_trajectory` and `optimize_racing_path_sequential`) are PROVIDED because:

- They're mostly boilerplate scipy code
- The real learning was in implementing cost/constraint functions (which you did!)
- You'll learn more by USING them and understanding the results

Think of it like this:

- **You built the engine** (cost functions, constraints)
- **We provided the chassis** (scipy wrapper code)
- **Now let's race!**

Ready? Let's optimize some trajectories!

✓ 7.3.1 Simple Path Optimization (No Hoops)

Let's start simple - optimizing a path from point A to B without hoops.

What to watch for:

- RRT path (angular, suboptimal)
- Optimized path (smooth, efficient)
- Cost reduction
- Constraint satisfaction

```
# Define simple scenario
start_simple = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(2, 2, 5))
goal_simple = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(8, 8, 5))

print("Step 1: Running RRT...")
rrt_simple, parents_simple = run_rrt(
    start_simple, goal_simple,
    generate_random_pose, steer, helpers.distance_between_poses,
    find_nearest_pose, threshold=2.0
)
path_simple = get_rrt_path(rrt_simple, parents_simple)
print(f"RRT path: {len(path_simple)} waypoints")

print("\nStep 2: Optimizing trajectory...")
optimized_simple, success_simple, info_simple = helpers.optimize_trajectory(
    rrt_path=path_simple,
    start_pose=start_simple,
    goal_position=goal_simple.translation(),
    hoops=[],
    obstacles=[],
    N=20,
    dt=0.1,
    weights={'thrust': 0.1, 'angular': 1.0, 'smoothness': 5.0},
    # Pass student-implemented functions (TODOs 17-28)
    initialize_from_rrt_robust=initialize_from_rrt_robust,
    dynamics_constraints_robust=dynamics_constraints_robust,
    boundary_constraints_robust=boundary_constraints_robust,
    collision_constraints_optimized=collision_constraints_optimized,
    cost_function_tuned=cost_function_tuned,
    unpack_decision_vars=unpack_decision_vars
)

if success_simple:
    print(f"\n✅ SUCCESS! Cost: {info_simple['cost']:.2f}, Iterations: {info_simple['iterations']}")")
else:
    print(f"\n❌ Optimization failed, using RRT path")

print("\nStep 3: Visualizing...")

# Visualize comparison
fig = helpers.visualize_rrt_vs_optimized_comparison(
    path_simple, optimized_simple, start_simple, goal_simple,
    title="Demo 1: RRT vs Optimized"
)
fig.show()
```

```
Step 1: Running RRT...
RRT path: 8 waypoints
```

Step 2: Optimizing trajectory...

```
=====
TRAJECTORY OPTIMIZATION
=====
Knot points: 20, Time step: 0.1s, Duration: 2.0s
Decision variables: 206
=====
```

Initializing from RRT...
Setting up constraints...

Optimizing with SLSQP...

```
/usr/local/lib/python3.11/dist-packages/scipy/optimize/_slsqp_py.py:435: RuntimeWarning:
```

Values in x were outside bounds during a minimize step, clipping to bounds

```
Optimization terminated successfully  (Exit mode 0)
    Current function value: 0.0015510104179032903
    Iterations: 122
    Function evaluations: 25320
    Gradient evaluations: 122
```

```
=====
✓ OPTIMIZATION SUCCEEDED
=====
Status: Optimization terminated successfully
Iterations: 122
Final cost: 0.0016
Optimization time: 32.18s
Max constraint violation: 0.000001
=====
```

✓ SUCCESS! Cost: 0.00, Iterations: 122

Step 3: Visualizing...

Demo 1: RRT vs Optimized

What Just Happened?

- **RRT path (cyan)**: Angular, explores randomly
- **Optimized path (magenta)**: Smooth, nearly straight
- **Cost reduction**: The optimizer minimized thrust, angular velocity, and jerk
- **Physics respected**: All dynamics constraints satisfied

This is trajectory optimization in action! 🤘

7.3.2 Racing with Hoops (No Obstacles)

Now let's tackle the real challenge - navigating through 4 hoops! We will first redo the RRT path from drone racing and then optimize this path for comparison!

```
# Get racing setup
hoops_demo2 = helpers.get_hoops()
targets_demo2 = helpers.get_targets()
start_demo2 = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(1, 3, 8))
```

```
print("Step 1: Running RRT through hoops...")
rrt_race_path = drone_racing_rrt(start_demo2, targets_demo2)
print(f"RRT path: {len(rrt_race_path)} waypoints")

print("\nStep 2: Optimizing racing trajectory...")
optimized_race, success_race, info_race = helpers.optimize_racing_path_sequential(
    rrt_path=rrt_race_path,
    start_pose=start_demo2,
    hoops=hoops_demo2,
    obstacles=[],
    N=25,
    dt=0.1,
    weights={'thrust': 0.05, 'angular': 0.5, 'smoothness': 5.0},
    # Pass student-implemented functions (TODOS 17-28)
    initialize_from_rrt_robust=initialize_from_rrt_robust,
    dynamics_constraints_robust=dynamics_constraints_robust,
    boundary_constraints_robust=boundary_constraints_robust,
    collision_constraints_optimized=collision_constraints_optimized,
    cost_function_tuned=cost_function_tuned,
    unpack_decision_vars=unpack_decision_vars
)

if success_race:
    print(f"\n✅ Racing optimization SUCCESS!")
else:
    print(f"\n⚠ Some segments may have failed")

print("\nStep 3: Visualizing...")

# Visualize racing comparison
fig = helpers.drone_racing_path_comparison(
    hoops_demo2, start_demo2, rrt_race_path, optimized_race,
    title="Demo 2: RRT vs Optimized Racing"
)
fig.show()
```

```
Step 1: Running RRT through hoops...
RRT path: 28 waypoints
```

```
Step 2: Optimizing racing trajectory...
```

```
=====
SEQUENTIAL RACING PATH OPTIMIZATION
=====
Hoops: 4, Knot points per segment: 25
=====
```

```
Splitting path into segments:
Hoop 1 at RRT waypoint 5
Hoop 2 at RRT waypoint 13
Hoop 3 at RRT waypoint 19
Hoop 4 at RRT waypoint 26
```

```
-----
SEGMENT 1/4 → Hoop 1
-----
```

```
=====
TRAJECTORY OPTIMIZATION
=====
```

```
Knot points: 25, Time step: 0.1s, Duration: 2.5s
Decision variables: 256
=====
```

```
Initializing from RRT...
Setting up constraints...
```

```
Optimizing with SLSQP...
```

```
Optimization terminated successfully (Exit mode 0)
  Current function value: 0.0005225140006579233
  Iterations: 270
  Function evaluations: 69473
  Gradient evaluations: 270
```

```
=====
✓ OPTIMIZATION SUCCEEDED
=====
```

```
Status: Optimization terminated successfully
Iterations: 270
Final cost: 0.0005
Optimization time: 125.05s
Max constraint violation: 0.000000
=====
```

```
✓ Segment 1 succeeded
```

```
-----
SEGMENT 2/4 → Hoop 2
-----
```

```
=====
Racing Analysis
=====
```

```
Knot points: 25, Time step: 0.1s, Duration: 2.5s
```

RRT (cyan): Gets through hoops but with sharp turns **Optimized (magenta)**: Smooth arcing turns, consistent velocity

The optimized path is:

- Initializing from RRT...
- Setting up constraints...
- More efficient (better battery life)
- Optimizing with SLSQP...
- Easier to track (predictable for controllers)

Optimization terminated successfully (Exit mode 0)
Real drone racing teams use exactly this approach!

```
Current function value: 0.010625777037038486
Iterations: 271
Function evaluations: 69745
```

7.3.3 Velocity & Acceleration Profiles

Let's look "under the hood" at the velocity and acceleration.

```
=====
if success_race and 'states' in info_race:
    states_opt = info_race['states']
    controls_opt = info_race['controls']

    # Compute velocities
```

```

dt = 0.1
velocities = np.diff(states_opt[:, :3], axis=0) / dt
accelerations = np.diff(velocities, axis=0) / dt

speeds = np.linalg.norm(velocities, axis=1)
print(f"Velocity stats:")
print(f"  Max speed: {speeds.max():.2f} m/s")
print(f"  Mean speed: {speeds.mean():.2f} m/s")

accel_mags = np.linalg.norm(accelerations, axis=1)
print(f"\nAcceleration stats:")
print(f"  Max: {accel_mags.max():.2f} m/s²")
print(f"  Mean: {accel_mags.mean():.2f} m/s²")

print(f"\nControl stats:")
print(f"  Thrust range: [{controls_opt[:, 3].min():.1f}, {controls_opt[:, 3].max():.1f}]")
print(f"  Mean thrust: {controls_opt[:, 3].mean():.1f} (hover=10.0)")

# Plot
fig = helpers.plot_velocity_acceleration_profiles(velocities, accelerations, controls_opt, dt)
fig.show()
else:
    print("⚠ Profile data not available")

```

OPTIMIZATION SUCCEEDED
Profile data not available

Status: Optimization terminated successfully

Iterations: 249

Optimization time: 112.61s

Max constraint violation: 0.000000
=====

7.3.4 Path Metrics Comparison

Let's quantify the improvement on how our optimized path is better than the non-optimized one!

```

def compute_path_length(path):
    """Compute total path length"""
    length = 0.0
    for i in range(len(path)-1):
        pos1 = np.array(path[i].translation())
        pos2 = np.array(path[i+1].translation())
        length += np.linalg.norm(pos2 - pos1)
    return length

```

Knot points: 25, Time step: 0.1s, Duration: 2.5s

Decision variables: 256

PART 8: The Final Challenge: Racing with Obstacles!

This is it. Everything you've built - RRT planning, terminal velocity dynamics, trajectory optimization, constraint handling - all comes together RIGHT NOW.

Optimizing with SLSQP...

You've seen your drone fly smooth paths. You've watched it race through hoops. But now? Now we add **obstacles** to the mix. This is the ultimate test of your optimizer: successfully (Exit mode 0)

Current function value: 0.0011222463814686916

Iterations: 216

Function evaluations: 55596

Gradient evaluations: 216

1. Navigate through 4 aerial hoops (you've done this before!)

2. Dodge obstacles placed strategically in your path (NEW!)

3. Satisfy All constraints simultaneously:

Status: Optimization terminated successfully

Iterations: 198 Dynamics constraints (physics must be obeyed)

Final cost: 0.011

Optimization time: 97.49s

Max % Hoop constraints (pass through centers)

Collision constraints (avoid obstacles with safety margin)=

4. Optimize for smoothness while doing all of the above

Why This Matters:

Concatenating segments...

This isn't just a demo anymore. This is **real autonomous drone racing**. The kind of problem that gets drones through disaster zones, warehouse navigation, and yes - actual competitive racing leagues.

Your RRT will plan a collision-free path. Your optimizer will make it fast and smooth. Your constraints will keep it safe and accurate.

Racing optimization SUCCESS!

The Stakes:

Step 3: Visualizing...

- If your dynamics constraints fail -> your drone violates physics

- If your boundary constraints fail you miss the start or goal
- If your collision constraints fail -> you crash into obstacles
- If your optimization succeeds → **YOU WIN** 🎉

Take a deep breath. Run the cells below. Watch your drone navigate the obstacle course like a boss.

Let's do this! 🚀

```
print("\n" + "="*80)
print("※※ FINAL CHALLENGE: RACING WITH OBSTACLES (EASY)")
print("=".join(["="]*80))

# Setup
hoops_final = helpers.get_hoops()
targets_final = helpers.get_targets()
start_final = gtsam.Pose3(gtsam.Rot3(), gtsam.Point3(1, 3, 8))
obstacles_easy = helpers.get_obstacles_easy()

print(f"Configuration:")
print(f"  - Start: {start_final.translation()}")
print(f"  - Hoops: {len(hoops_final)}")
print(f"  - Obstacles: {len(obstacles_easy)} (EASY)")

print("\n" + "-"*80)
print("STAGE 1: RRT WITH OBSTACLES")
print("-" * 80)

import time
start_time = time.time()
rrt_final_path = drone_racing_rrt_with_obstacles(start_final, targets_final, obstacles_easy)
rrt_time = time.time() - start_time

print(f"RRT completed in {rrt_time:.2f}s")
print(f"Path: {len(rrt_final_path)} waypoints")

has_collision, _ = helpers.check_path_collision(rrt_final_path, obstacles_easy)
if not has_collision:
    print(f"RRT is Collision-free!")
else:
    print(f"RRT may have collisions")
```

```
=====
※※ FINAL CHALLENGE: RACING WITH OBSTACLES (EASY)
=====

Configuration:
  - Start: [1. 3. 8.]
  - Hoops: 4
  - Obstacles: 2 (EASY)

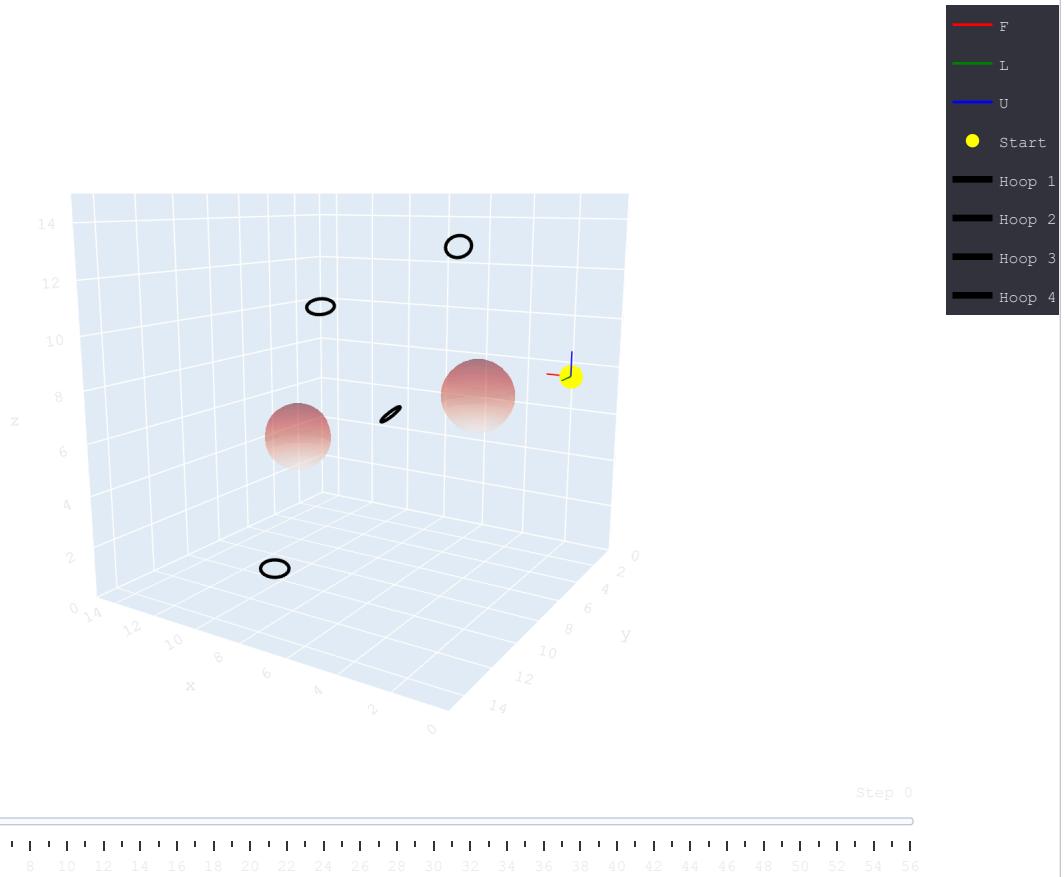
-----
STAGE 1: RRT WITH OBSTACLES
-----
Planning to hoop 0...
  Using best path (distance 0.56 from hoop)
Planning to hoop 1...
  Using best path (distance 1.59 from hoop)
Planning to hoop 2...
  Using best path (distance 1.86 from hoop)
Planning to hoop 3...
  Using best path (distance 1.62 from hoop)
RRT completed in 9.27s
Path: 57 waypoints
RRT is Collision-free!
```

```
# Visualize RRT racing path with obstacles once again
helpers.drone_racing_path_with_obstacles(
    hoops=helpers.get_hoops(),
    start=start_final,
    path=rrt_final_path,
    obstacles=obstacles_easy
)

print(f"\nTotal racing path length: {len(rrt_final_path)} waypoints")

# Verify collision-free
has_collision, _ = helpers.check_path_collision(rrt_final_path, obstacles_easy)
```

```
if not has_collision:
    print(f"RRT is Collision-free!")
else:
    print(f"RRT may have collisions")
```



Total racing path length: 57 waypoints
RRT is Collision-free!

```
print("\n" + "-"*80)
print("STAGE 2: TRAJECTORY OPTIMIZATION")
print("-"*80)

opt_start = time.time()
optimized_final, success_final, info_final = helpers.optimize_racing_path_sequential(
    rrt_path=rrt_final_path,
    start_pose=start_final,
    hoops=hoops_final,
    obstacles=obstacles_easy,
    N=25,
    dt=0.1,
    weights={'thrust': 0.05, 'angular': 0.5, 'smoothness': 5.0},
    initialize_from_rrt_robust=initialize_from_rrt_robust,
    dynamics_constraints_robust=dynamics_constraints_robust,
    boundary_constraints_robust=boundary_constraints_robust,
    collision_constraints_optimized=collision_constraints_optimized,
    cost_function_tuned=cost_function_tuned,
    unpack_decision_vars=unpack_decision_vars
)
opt_time = time.time() - opt_start

if success_final:
    print(f"\n✅ OPTIMIZATION SUCCESS!")
    print(f"  Time: {opt_time:.2f}s")
    print(f"  Total time: {rrt_time + opt_time:.2f}s")

    has_collision_opt, _ = helpers.check_path_collision(optimized_final, obstacles_easy)
    if not has_collision_opt:
        print(f"Optimized path is collision-free!")
    print(f"Passes through all {len(hoops_final)} hoops")
```

```

else:
    print(f"\n Optimization had issues, using RRT path")
    optimized_final = rrt_final_path

-----
STAGE 2: TRAJECTORY OPTIMIZATION
-----

=====
SEQUENTIAL RACING PATH OPTIMIZATION
=====
Hoops: 4, Knot points per segment: 25
=====

Splitting path into segments:
Hoop 1 at RRT waypoint 8
Hoop 2 at RRT waypoint 20
Hoop 3 at RRT waypoint 36
Hoop 4 at RRT waypoint 54

-----
SEGMENT 1/4 → Hoop 1
-----

=====
TRAJECTORY OPTIMIZATION
=====
Knot points: 25, Time step: 0.1s, Duration: 2.5s
Decision variables: 256
=====

Initializing from RRT...
Setting up constraints...
    Added collision avoidance for 2 obstacles

Optimizing with SLSQP...
/usr/local/lib/python3.11/dist-packages/scipy/optimize/_slsqp.py:435: RuntimeWarning:
Values in x were outside bounds during a minimize step, clipping to bounds

Optimization terminated successfully      (Exit mode 0)
    Current function value: 0.0008160071503188936
    Iterations: 234
    Function evaluations: 60218
    Gradient evaluations: 234

=====
✓ OPTIMIZATION SUCCEEDED
=====
Status: Optimization terminated successfully
Iterations: 234
Final cost: 0.0008
Optimization time: 118.01s
Max constraint violation: 0.000007
=====

✓ Segment 1 succeeded
-----
```

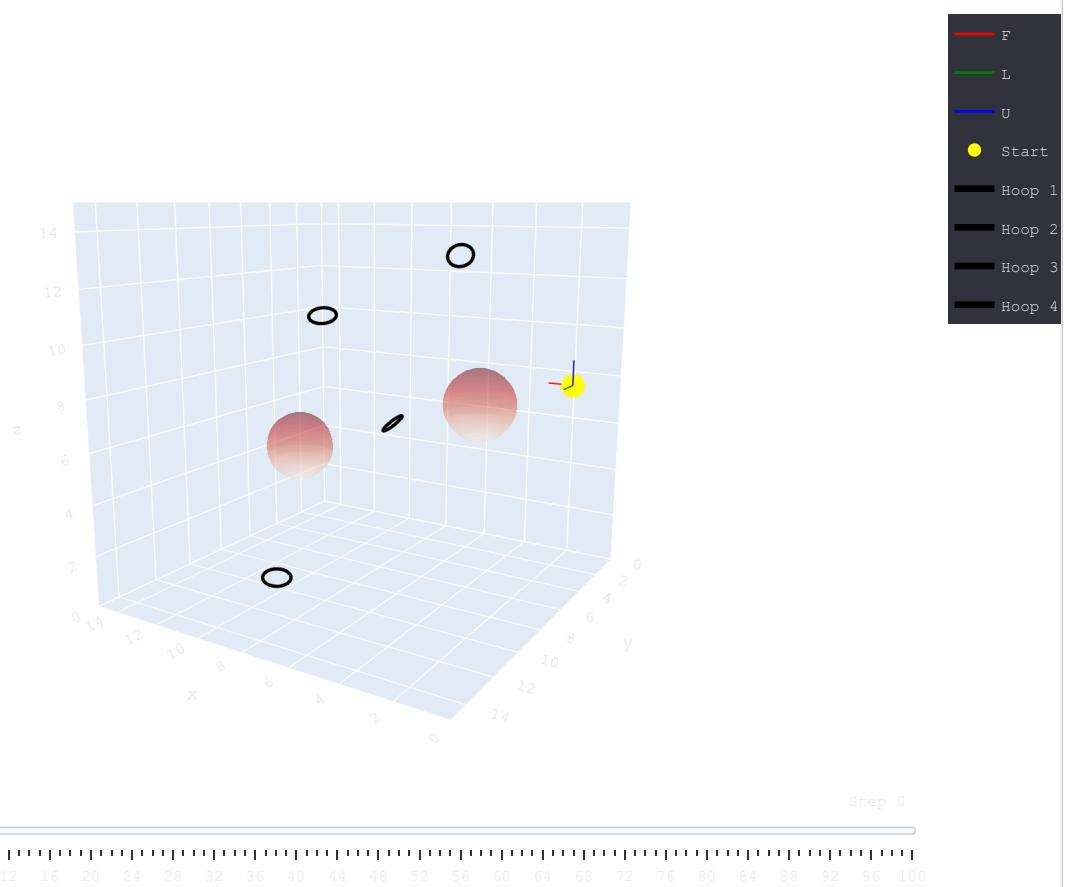
```

print("\n" + "-"*80)
print("STAGE 3: VISUALIZATION")
print("-"*80)

helpers.drone_racing_path_with_obstacles(
    hoops_final, start_final, optimized_final, obstacles_easy,
)

print("\n👉 Interactive 3D visualization above!")
print("    - Rotate: mouse drag")
print("    - Zoom: scroll wheel")
print("    - Pan: right-click drag")
```

STAGE 3: VISUALIZATION



- 💡 Interactive 3D visualization above!
- Rotate: mouse drag
 - Zoom: scroll wheel
 - Pan: right-click drag

```
# Path Metric Evaluation:
length_rrt_final = compute_path_length(rrt_final_path)
length_opt_final = compute_path_length(optimized_final)

print("PATH COMPARISON:")
print(f" RRT path length: {length_rrt_final:.2f} m")
print(f" Optimized path length: {length_opt_final:.2f} m")
print(f" Reduction: {(length_rrt_final-length_opt_final)/length_rrt_final*100:.1f}%")
print(f"\n RRT waypoints: {len(rrt_final_path)}")
print(f" Optimized waypoints: {len(optimized_final)}")

if success_final:
    print(f"\n💡 The optimized path is smoother and more efficient!")
```

```
PATH COMPARISON:
RRT path length: 85.77 m
Optimized path length: 41.21 m
Reduction: 52.0%

RRT waypoints: 57
Optimized waypoints: 101
```

💡 The optimized path is smoother and more efficient!

Final Reflection Questions:

Part A: Optimization Sensitivity Analysis

Similar to how we analyzed drone dynamics earlier, let's analyze how our optimization hyperparameters affect the final trajectory. Consider the following scenarios for the trajectory optimization. For each case, hypothesize (or test!) what will happen to the **path shape**, **speed**, and **computation time**.

Scenario	w_smoothness	w_thrust	N (Knot Points)	Description
1. Baseline	5.0	0.05	25	Balanced parameters
2. Ultra Smooth	50.0	0.05	25	Very high smoothness penalty
3. Aggressive	0.1	1.0	25	Low smoothness, high thrust cost
4. Low Res	5.0	0.05	10	Very few knot points
5. High Res	5.0	0.05	100	Many knot points

Questions

- Smoothness vs. Feasibility:** In Scenario 2 (Ultra Smooth), what happens to the drone's ability to pass through narrow hoops or avoid obstacles? Does the path become "too" simple?
- Aggressiveness:** In Scenario 3, does the drone fly faster or slower? Does it cut corners more tightly? What happens to the control inputs (thrust/body rates)?
- Resolution Trade-off:** Compare Scenarios 4 and 5. How does N affect the *validity* of the trajectory (collision checking resolution) vs. the *optimality* and *computation time*? Why can't we just use N=1000 for everything?

Part B: Overall Reflection

- RRT vs. Optimization:** How do these two approaches complement each other? Why is RRT good for initialization but bad for final execution? Why is Optimization good for smoothing but bad for finding a path from scratch?
- The "Reality Gap":** You implemented a physics-based model ($F = ma$, drag, etc.). In the real world, what other factors (wind, battery voltage sag, sensor noise, communication delay) would make this difficult? How would you modify your controller to handle them?
- Course Feedback:** What was the most "Aha!" moment for you in this assignment?

▼ Response:

PART A:

Response:

Question 1 - Ultra Smooth (w_smoothness = 50.0):

- The path likely becomes too smooth to the point it cannot make tight maneuvers
- Might have difficulties passing through tight hoop passages or narrow gaps between obstacles
- High smoothness penalty essentially forces the optimizer to avoid any sharp turns
- Could end up taking longer detours just to keep everything gentle
- The path could be "too simple" and miss out on optimal shortcuts that require sharper movements
- Basically trading performance for comfort

Question 2 - Aggressive (w_smoothness = 0.1, w_thrust = 1.0):

- Drone would actually fly more slowly, probably, since high thrust cost discourages the use of maximum power
- Might take sharper corners since smoothness penalty is low - less concerned about jerkiness
- Control inputs would be more erratic: sudden changes in thrust and orientation
- Low smoothness weight means that it can make aggressive maneuvers without penalty
- But the high thrust cost means it tries to minimize energy use, so it won't go at full speed
- Probably weird combination of sharp turns but conservative power usage

Question 3 - Resolution (N=10 vs N=100):

- Low N (10 knots): Faster computation, but really coarse path. Might miss obstacles between waypoints or violate dynamics constraints
- High N (100 knots): Super detailed path, catches all collisions, but takes forever to optimize and might overfit
- Can't use N=1000 because optimizer would have like 10,000 decision variables - would take hours to converge
- Also, more knot points mean more constraints to satisfy, making the problem way harder
- There's a sweet spot around N=20-30 where you get enough detail without killing performance
- Low N might look smooth in the optimizer but jerky when actually executed
- High N is more accurate but computational cost explodes - not practical for real-time planning

PART B:

Response:

Question 1 - RRT vs. Optimization:

- RRT is great at finding a path that avoids obstacles - it explores the space randomly until it finds something that works
- But RRT paths are super ugly - lots of random turns, not smooth at all
- Optimization is good at making paths smooth and efficient, that minimize energy and jerk
- But optimization from scratch would struggle to avoid obstacles: might get stuck in local minima, or violate collision constraints
- Using them together is perfect: RRT gives you a feasible starting point that's collision-free, then optimization polishes it

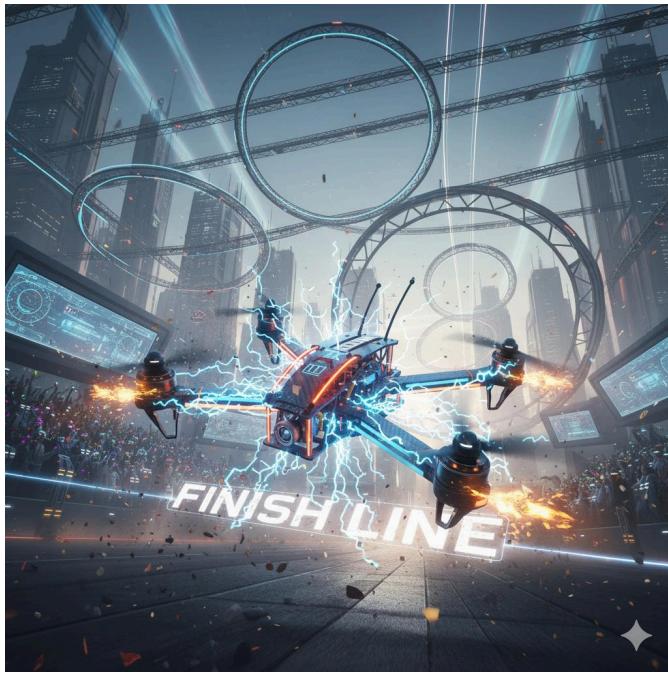
Question 2 - Reality Gap:

- Wind would mess everything up, our terminal velocity model assumes calm air
- Battery voltage decreases as it discharges, therefore, the thrust characteristics change during a flight
- Sensor noise means we don't actually know our exact position/orientation
- Delay in communication from commanding to motors responding
- Motor dynamics are not instantaneous. There is spinup time and momentum
- Props can stall if you pitch too aggressively
- GPS drift in outdoor environments, or camera tracking errors indoors
- Temperature influences air density and performance of motor
- The model is deterministic, but the real world is super noisy and unpredictable

Question 3 - Course Feedback:

- The largest "aha" moment was in actually seeing real trajectory optimization work, with all the constraints
- Things like realizing that you want equality constraints (physics, boundaries) AND inequality constraints (obstacles)
- Also realising why just doing RRT isn't enough - you actually need optimization to make it flyable
- The terminal velocity stuff was cool: seeing how thrust and orientation combine to create movement
- Honestly, the whole thing coming together toward the end when the drone could smoothly fly through hoops while avoiding obstacles was satisfying
- Made me appreciate how complex autonomous drone flight actually is compared to just "point and go"

🎉 Victory!



This is the end of the assignment. Everything after this point is extra credit. (Extra credit will be updated shortly)

Your drone just navigated through hoops AND dodged obstacles. Let that sink in for a second.

This is the kind of tech that:

- Powers Amazon delivery drones navigating urban environments
- Enables search-and-rescue drones in disaster zones
- Wins actual drone racing competitions (DRL, MultiGP)

- Gets used in Hollywood for autonomous aerial cinematography
 - Lands rovers on Mars (okay, different dynamics, but same optimization principles!)
-

💪 What You Can Do Now:

- **Impress recruiters:** "I implemented nonlinear trajectory optimization with collision avoidance for autonomous drones"
 - **Side projects:** Hook this up to a real drone (Tello EDU is like \$100)
 - **Research:** Extend it to multi-drone coordination, time-optimal trajectories, or learning-based planning
 - **Competitions:** Seriously, there are drone racing leagues that need exactly this kind of work but with better planners!
-

🔥 Hungry for More?

If you're the kind of person who thinks "that was too easy" (it wasn't, but we respect the confidence), scroll down.

The EXTRA CREDIT challenge below will separate the good from the legendary. 🙌

▼ ⚡ EXTRA CREDIT: The Gauntlet

The Challenge That Breaks Most Students

The easy course? That was a tutorial. **This is the main level.**

What You're Up Against:

The HARD Obstacle Course:

- **8 obstacles** (4x more than before)
- **Narrow gaps** between obstacles (your safety margin will be tested)
- **Complex 3D geometry** (boxes AND spheres, positioned to trap bad planners)
- **Longer computation** (RRT might need multiple attempts)
- **Tighter optimization** (SLSQP might fail, trust-constr will earn its keep)

You Have Total Freedom:

Unlike the main assignment, **you can use ANY path planning algorithm you want:**

- Stick with RRT (the classic)
- Implement RRT* (asymptotically optimal)
- Try RRT-Connect (bidirectional search)
- Use informed RRT* (if you're showing off)
- Implement MPPI, A*, DWA (if you really have that much time!)
- Design a custom heuristic planner
- Combine multiple strategies
- **OR Surprise us!**

The only rule: **Don't hardcode the solution.** Your planner must work for arbitrary obstacle configurations.

⚠ The Box Obstacle Twist:

Here's the kicker: **The easy course only had spherical obstacles.** Your `collision_constraints_optimized` function (TODO 26) probably only handles sphere collision checks.

The hard course has BOX obstacles too.

To complete this challenge, you need to:

1. **Extend** `collision_constraints_optimized` to handle both sphere AND box obstacles
2. **Implement box collision constraints** (hint: think about axis-aligned bounding boxes—check if drone position is within [min_corner, max_corner] plus safety margin)
3. **Make it work seamlessly** with your existing optimizer

Define the `collision_constraints_optimized` or any new functions again which you want to modify from the previous sections (don't modify them in place as it can hinder your assignment grades)

🏆 Grading Rubric (Worth It):

Core Challenge (Multiple Obstacles):

- +10% Implement box collision constraints accurately
- +5% Optimization converges (constraints satisfied)
- +10% Final optimized path avoids ALL obstacles (spheres + boxes)

Bonus Points (Flex Zone):

- +20% Use a non-RRT planner successfully (RRT*, MPPI, etc.) which avoids all obstacles
- +5% Document your approach with detailed comments/markdown

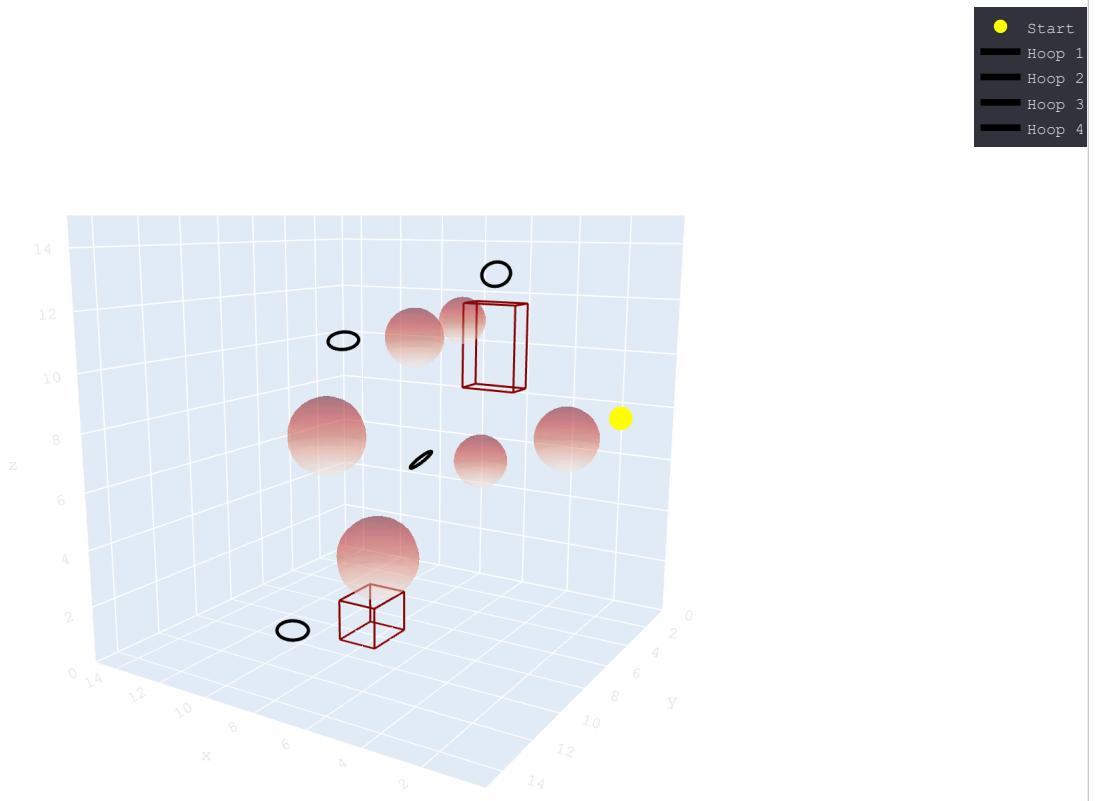
Total Possible: 50% of extra credit value

Extra Credit Grading Scale:

- **Assignment base:** 100 points
- **Extra credit raw:** Up to 50 points (percentages above)
- **How it works:** Your extra credit percentage is scaled proportionally to contribute up to bonus of upto half of the total assignment marks

Bottom line: Every bit counts, and perfection here can seriously boost your grade!

```
# Before Commiting for extra credit, let's see how the whole Racing environment looks like with 8 obstacles added instead of 2
start_race = gtsam.Pose3(r=gtsam.Rot3.Yaw(math.radians(45)), t=gtsam.Point3(1, 3, 8))
helpers.drone_racing_path_with_obstacles(helpers.get_hoops(), start_race, [], obstacles=helpers.get_obstacles_hard())
```



⌄ 🧠 Why This Is Hard:

- **Box obstacles are NEW** -> Your collision function only handles spheres right now
- **RRT might fail** -> You'll need retry logic or adaptive parameters
- **Optimization might fail** -> You'll need good initialization
- **Collisions are sneaky** -> Subsampling must be dense enough, especially for boxes
- **Constraints conflict** -> Hoops might be near obstacles
- **Mixed geometry** -> Handling both sphere AND box constraints simultaneously
- **Computation time** -> You might wait 2-5 minutes per attempt

 Hints (Use Wisely):

1. **Box collision math:** For a point to be OUTSIDE a box, it must be outside on at least one face. Constraint per box: `max(0, -(pos[0] - box.max[0] - safety), -(box.min[0] - pos[0] - safety), ...)`
2. **RRT struggles?** Try increasing max iterations or adding goal bias
3. **Optimization fails?** Use more knot points (N=40-50) or adjust weights
4. **Collisions after optimization?** Check your `collision_constraints_optimized` function handles BOTH obstacle types
5. **Getting desperate?** The provided code has retry logic for RRT (see below)

 The Mindset:

This is **research-level** robotics. Real autonomous systems fail sometimes. Your job is to:

1. **Understand why** things fail
2. **Adapt your strategy**
3. **Persist until it works**

If you get this working, you're in the top 5% of students. If you get bonus points, you're ready for graduate-level robotics.

No pressure. Just glory. 

```
##### Student implements Custom planner and optimization here #####
# =====#
# EXTRA CREDIT: PRM - USES HOOPS DIRECTLY
# =====#
# =====#
import numpy as np
from typing import List, Dict, Tuple
import heapq

# =====#
# Extended Collision Constraints
# =====#

def collision_constraints_optimized_extended(z: np.ndarray, N: int, obstacles: List,
                                               subsample: int = 2) -> np.ndarray:
    states, controls = unpack_decision_vars(z, N)
    violations = []
    safety_margin = 0.5

    for k in range(0, N + 1, subsample):
        px, py, pz = states[k, 0:3]
        pos = np.array([px, py, pz])

        for obs in obstacles:
            if isinstance(obs, helpers.SphereObstacle):
                dist = np.linalg.norm(pos - np.array(obs.center))
                violations.append(obs.radius + safety_margin - dist)

            elif isinstance(obs, helpers.BoxObstacle):
                min_corner = np.array(obs.min_corner)
                max_corner = np.array(obs.max_corner)

                dist_to_min_x = (min_corner[0] - safety_margin) - pos[0]
                dist_to_max_x = pos[0] - (max_corner[0] + safety_margin)
                dist_to_min_y = (min_corner[1] - safety_margin) - pos[1]
                dist_to_max_y = pos[1] - (max_corner[1] + safety_margin)
                dist_to_min_z = (min_corner[2] - safety_margin) - pos[2]
                dist_to_max_z = pos[2] - (max_corner[2] + safety_margin)

                max_penetration = max(dist_to_min_x, dist_to_max_x,
                                      dist_to_min_y, dist_to_max_y,
                                      dist_to_min_z, dist_to_max_z)
                violations.append(max_penetration)

    return np.array(violations)

# =====#
# Helper: Extract position from ANY type
# =====#
def get_position(obj):
    """Extract [x, y, z] from any type"""

```

```

if isinstance(obj, np.ndarray):
    return obj.flatten()[:3].astype(float)
if isinstance(obj, (list, tuple)):
    return np.array(obj[:3], dtype=float)
if hasattr(obj, 'x') and callable(getattr(obj, 'x')):
    try:
        return np.array([obj.x(), obj.y(), obj.z()], dtype=float)
    except:
        pass
if hasattr(obj, 'translation') and callable(getattr(obj, 'translation')):
    try:
        t = obj.translation()
        if isinstance(t, np.ndarray):
            return t.flatten()[:3].astype(float)
        elif hasattr(t, 'x') and callable(getattr(t, 'x')):
            return np.array([t.x(), t.y(), t.z()], dtype=float)
        else:
            return np.array(t, dtype=float).flatten()[:3]
    except:
        pass
# Try center attribute (for hoops)
if hasattr(obj, 'center'):
    c = obj.center
    if isinstance(c, np.ndarray):
        return c.flatten()[:3].astype(float)
    elif hasattr(c, 'x') and callable(getattr(c, 'x')):
        return np.array([c.x(), c.y(), c.z()], dtype=float)
    else:
        return np.array(c, dtype=float).flatten()[:3]
try:
    return np.array(obj, dtype=float).flatten()[:3]
except:
    raise ValueError(f"Cannot extract position from {type(obj)}")

def get_hoop_center(hoop):
    """Extract center position from a hoop object"""
    # Trying different attributes that hoops might have
    if hasattr(hoop, 'center'):
        c = hoop.center
        if callable(c):
            c = c()
        if isinstance(c, np.ndarray):
            return c.flatten()[:3].astype(float)
        elif hasattr(c, 'x'):
            if callable(getattr(c, 'x')):
                return np.array([c.x(), c.y(), c.z()], dtype=float)
            else:
                return np.array([c.x, c.y, c.z], dtype=float)
        return np.array(c, dtype=float).flatten()[:3]

    if hasattr(hoop, 'translation') and callable(getattr(hoop, 'translation')):
        t = hoop.translation()
        if isinstance(t, np.ndarray):
            return t.flatten()[:3].astype(float)
        elif hasattr(t, 'x') and callable(getattr(t, 'x')):
            return np.array([t.x(), t.y(), t.z()], dtype=float)

    if hasattr(hoop, 'x') and callable(getattr(hoop, 'x')):
        return np.array([hoop.x(), hoop.y(), hoop.z()], dtype=float)

    if isinstance(hoop, np.ndarray):
        return hoop.flatten()[:3].astype(float)

    return np.array(hoop, dtype=float).flatten()[:3]

# =====
# PRM Planner
# =====

class PRMPlanner:
    def __init__(self, bounds, n_samples=500, connection_radius=4.5, safety_margin=0.4):
        self.bounds = bounds
        self.n_samples = n_samples
        self.connection_radius = connection_radius
        self.safety_margin = safety_margin

```

```

self.nodes = []
self.edges = {}

def is_collision(self, pos, obstacles):
    for obs in obstacles:
        if isinstance(obs, helpers.SphereObstacle):
            if np.linalg.norm(pos - np.array(obs.center)) < obs.radius + self.safety_margin:
                return True
        elif isinstance(obs, helpers.BoxObstacle):
            min_c = np.array(obs.min_corner) - self.safety_margin
            max_c = np.array(obs.max_corner) + self.safety_margin
            if (min_c[0] <= pos[0] <= max_c[0] and
                min_c[1] <= pos[1] <= max_c[1] and
                min_c[2] <= pos[2] <= max_c[2]):
                return True
    return False

def is_edge_free(self, p1, p2, obstacles):
    dist = np.linalg.norm(p2 - p1)
    steps = max(int(dist / 0.3), 2)
    for i in range(steps + 1):
        point = p1 + (i / steps) * (p2 - p1)
        if self.is_collision(point, obstacles):
            return False
    return True

def build_roadmap(self, obstacles):
    print(" PRM: Building roadmap...")
    self.nodes = []
    self.edges = {}

    attempts = 0
    while len(self.nodes) < self.n_samples and attempts < self.n_samples * 20:
        attempts += 1
        point = np.array([
            np.random.uniform(self.bounds[0][0], self.bounds[0][1]),
            np.random.uniform(self.bounds[1][0], self.bounds[1][1]),
            np.random.uniform(self.bounds[2][0], self.bounds[2][1])
        ])
        if not self.is_collision(point, obstacles):
            idx = len(self.nodes)
            self.nodes.append(point)
            self.edges[idx] = []
            if len(self.nodes) % 100 == 0:
                print(f"     Sampled {len(self.nodes)}/{self.n_samples}...")

    print(f" PRM: Sampled {len(self.nodes)} nodes")

    connections = 0
    for i in range(len(self.nodes)):
        distances = []
        for j in range(len(self.nodes)):
            if i != j:
                dist = np.linalg.norm(self.nodes[i] - self.nodes[j])
                if dist < self.connection_radius:
                    distances.append((j, dist))

        distances.sort(key=lambda x: x[1])
        for j, dist in distances[:15]:
            if not any(n[0] == j for n in self.edges[i]):
                if self.is_edge_free(self.nodes[i], self.nodes[j], obstacles):
                    self.edges[i].append((j, dist))
                    self.edges[j].append((i, dist))
                    connections += 1

    print(f" PRM: Created {connections} edges")

def add_node(self, point, obstacles):
    idx = len(self.nodes)
    self.nodes.append(point.copy())
    self.edges[idx] = []

    for i in range(len(self.nodes) - 1):
        dist = np.linalg.norm(point - self.nodes[i])
        if dist < self.connection_radius * 1.5:
            if self.is_edge_free(point, self.nodes[i], obstacles):
                self.edges[idx].append((i, dist))

```

```

        self.edges[i].append((idx, dist))

    return idx

def dijkstra(self, start_idx, goal_idx):
    pq = [(0, start_idx)]
    came_from = {}
    cost = {start_idx: 0}

    while pq:
        curr_cost, curr = heapq.heappop(pq)

        if curr == goal_idx:
            path = []
            while curr in came_from:
                path.append(self.nodes[curr])
                curr = came_from[curr]
            path.reverse()
            return path

        for neighbor, edge_cost in self.edges.get(curr, []):
            new_cost = cost[curr] + edge_cost
            if neighbor not in cost or new_cost < cost[neighbor]:
                cost[neighbor] = new_cost
                came_from[neighbor] = curr
                heapq.heappush(pq, (new_cost, neighbor))

    return None

def plan(self, start_pos, goal_pos, obstacles):
    start_idx = self.add_node(start_pos, obstacles)
    goal_idx = self.add_node(goal_pos, obstacles)

    path = self.dijkstra(start_idx, goal_idx)

    if path:
        path.insert(0, start_pos.copy())
        path.append(goal_pos.copy())

    return path

# =====
# Drone Racing with PRM - USES HOOPS DIRECTLY
# =====

def drone_racing_prm(start, targets: List, obstacles: List = None,
                      hoops: List = None) -> List[gtsam.Pose3]:
    """
    Navigate through hoops using PRM.

    IMPORTANT: Uses 'hoops' parameter to get actual hoop centers!
    """

    if obstacles is None:
        obstacles = []

    bounds = [(0, 15), (0, 15), (0, 15)]
    drone_path = []

    current_pos = get_position(start)

    print("\n" + "*60)
    print("PRM (Probabilistic Roadmap) Planner")
    print("*60)

    # Build roadmap
    planner = PRMPlanner(bounds, n_samples=500, connection_radius=4.5)
    planner.build_roadmap(obstacles)

    print("\n" + "*60)
    print("Querying paths to each hoop...")
    print("*60)

    # Debugging: Printing what working with
    print(f"\n DEBUG: targets type = {type(targets)}")
    if len(targets) > 0:

```